CSCE 5300 Introduction to Big Data and Data Science

Lesson 5

Apache Lucene
Apache Solr

Overview

- Apache Lucene
- Apache Solr

Apache Lucene

Apache Lucene Highlights

- Fast, high performance, scalable search/IR library
- Open source
- Initially developed by Doug Cutting (Also author of Hadoop)
- Indexing and Searching
- Inverted Index of documents
- Provides advanced Search options like synonyms, stopwords, based on similarity, proximity.
- **http://lucene.apache.org/**

Lucene - Architecture

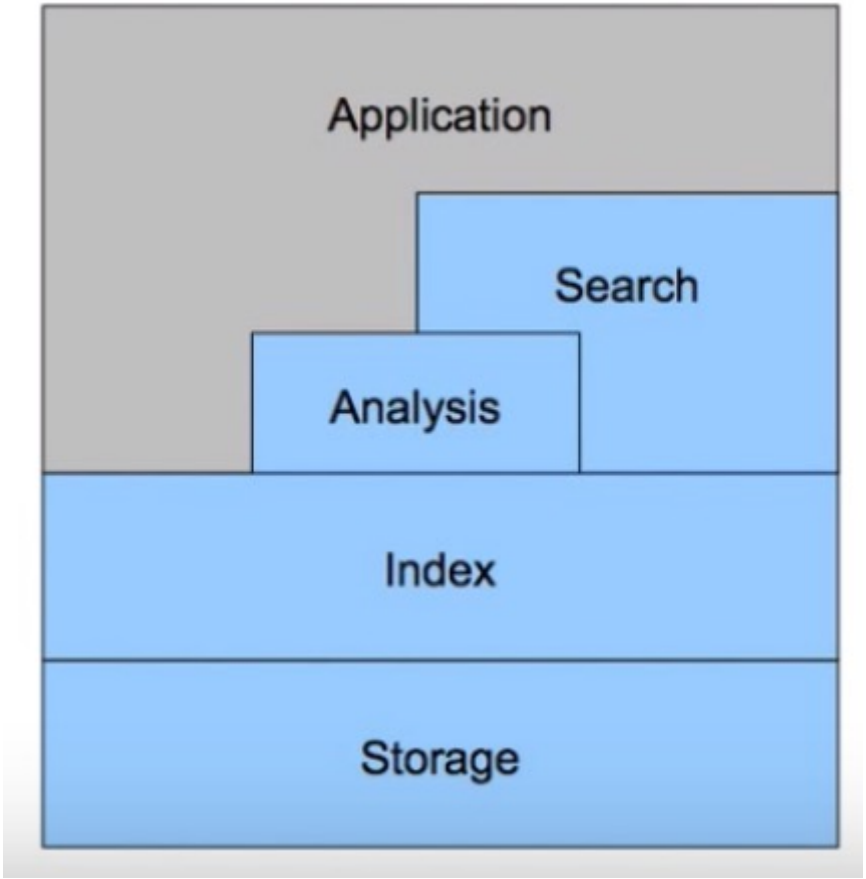

Lucene – Work Flow

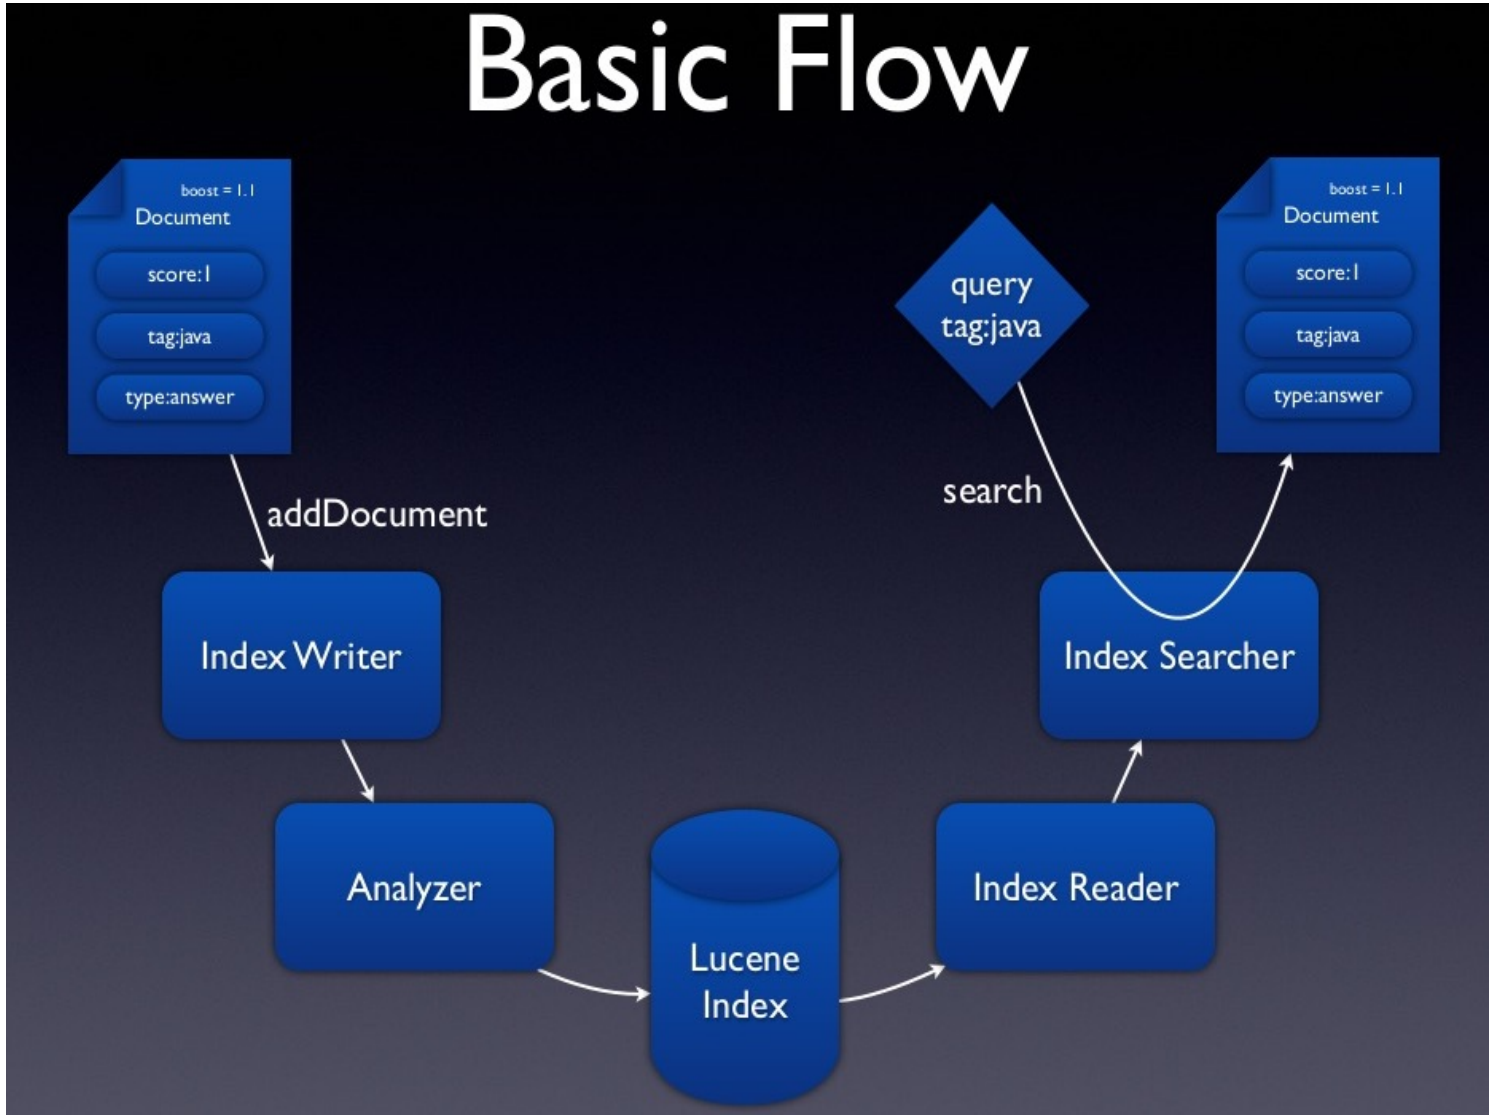

Lucene Internals - Inverted Index

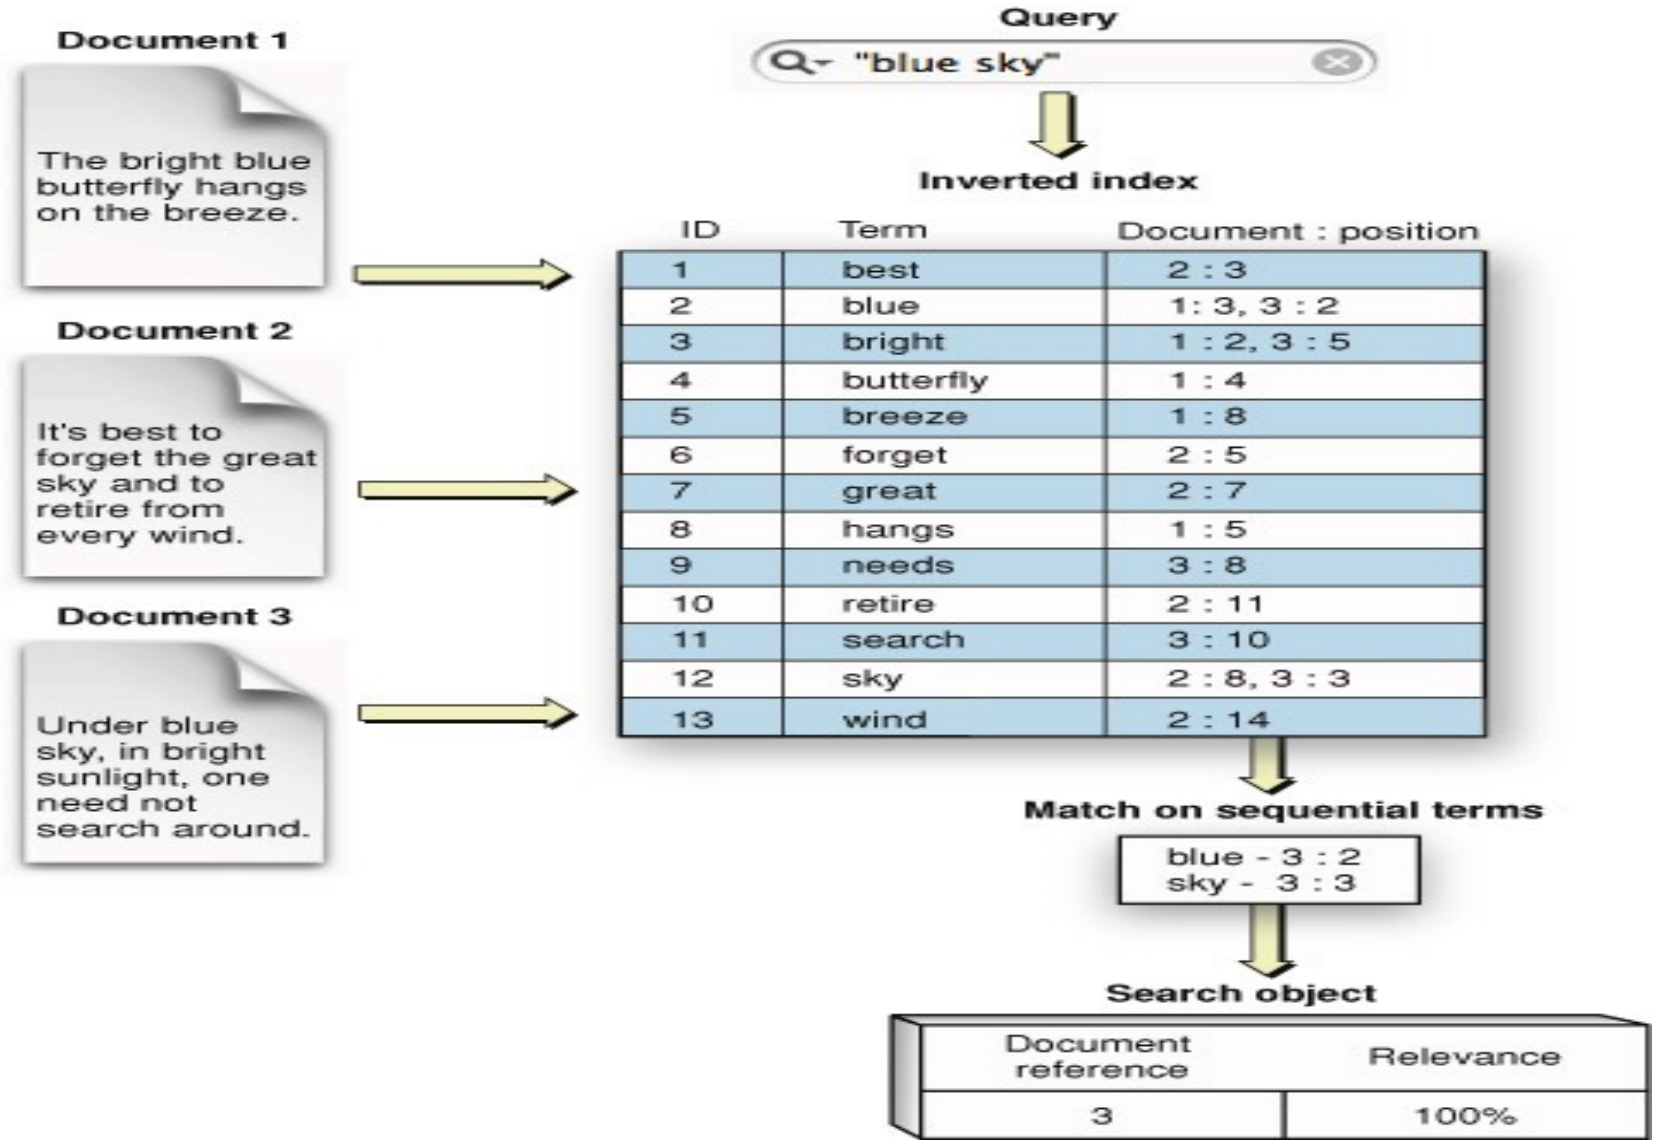

Credit: https://developer.apple.com/library/mac/documentation/userexperience/conceptual/SearchKitConcepts/searchKit_basics/searchKit_basics.html

Lucene - Indexing

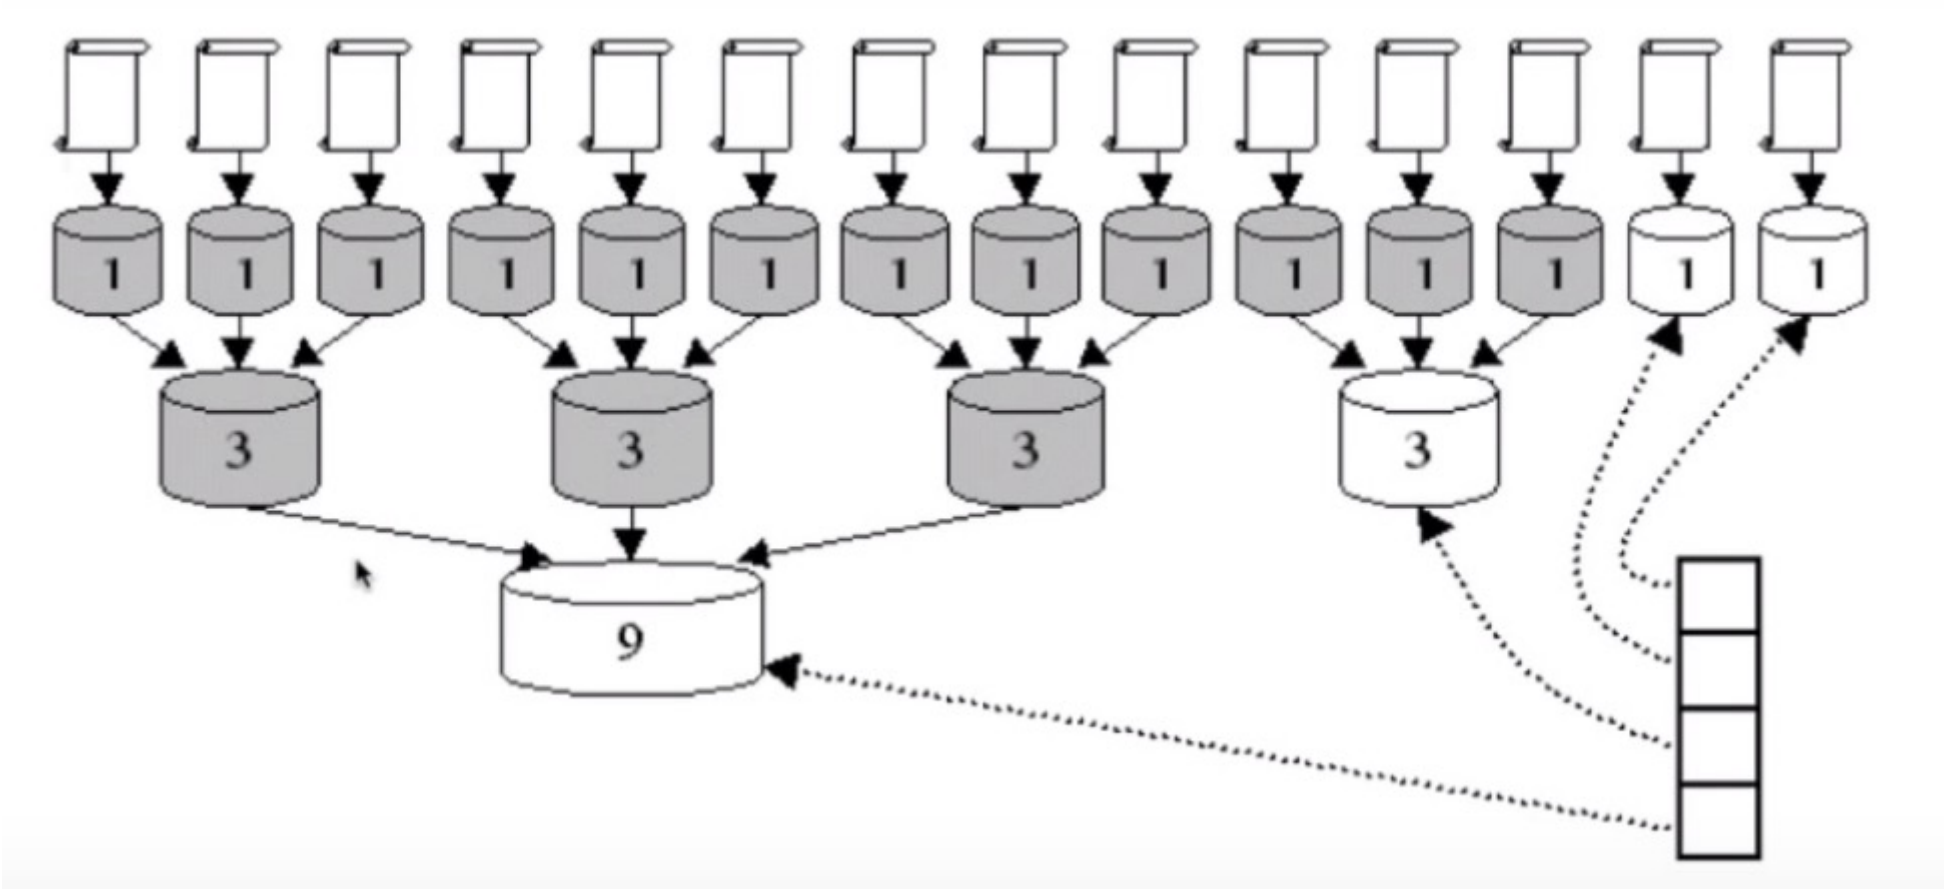

Indexing Pipeline

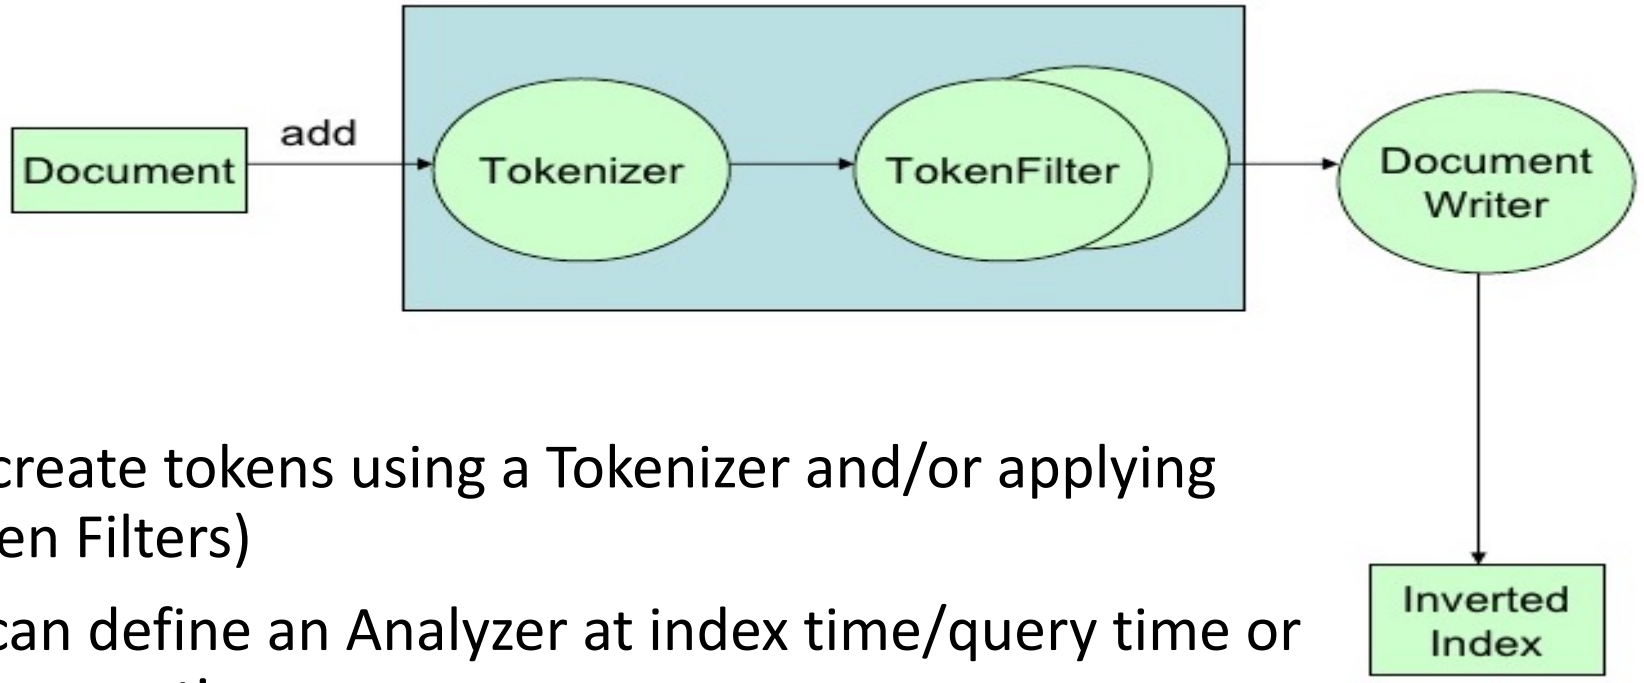

- Analyzer : create tokens using a Tokenizer and/or applying Filters (Token Filters)
- Each field can define an Analyzer at index time/query time or the both at same time.

Credit : http://www.slideshare.net/otisg/lucene-introduction

Analysis Process - Tokenizer

WhitespaceAnalyzer

Simplest built-in analyzer

The quick brown fox jumps over the lazy dog.

[The] [quick] [brown] [fox] [jumps] [over] [the] [lazy] [dog.]

Tokens

Analysis Process - Tokenizer

SimpleAnalyzer

Lowercases, split at non-letter boundaries

The quick brown fox jumps over the lazy dog.

[The] [quick] [brown] [fox] [jumps] [over] [the] [lazy] [dog.]

Tokens

Some common analyzer

- **WhitespaceAnalyzer** : Splits text at whitespaces, just as the name indicates. In fact, this is the only thing this analyzer does.
- **SimpleAnalyzer** : Splits text at non-letter characters and lowercases resulting tokens.
- **StopAnalyzer** : Splits text at non-letter characters, lowercases resulting tokens, and removes stopwords.
- **StandardAnalyzer** : Splits text using a grammar-based tokenization, normalizes and lowercases tokens, removes stopwords, and discards punctuations. It can be used to extract company names, e-mail addresses, model numbers, and so on. This analyzer is great for general usage.
- **SnowballAnalyzer**: This analyzer is similar to StandardAnalyzer with an additional SnowballFilter for stemming.

Apache Solr

Apache Solr

- Created by Yonik Seeley for CNET
- Enterprise Search platform for Apache Lucene
- Open source
- Highly reliable, scalable, fault tolerant
- Support distributed Indexing (SolrCloud), Replication, and load balanced querying
- **http://lucene.apache.org/solr**

High level overview

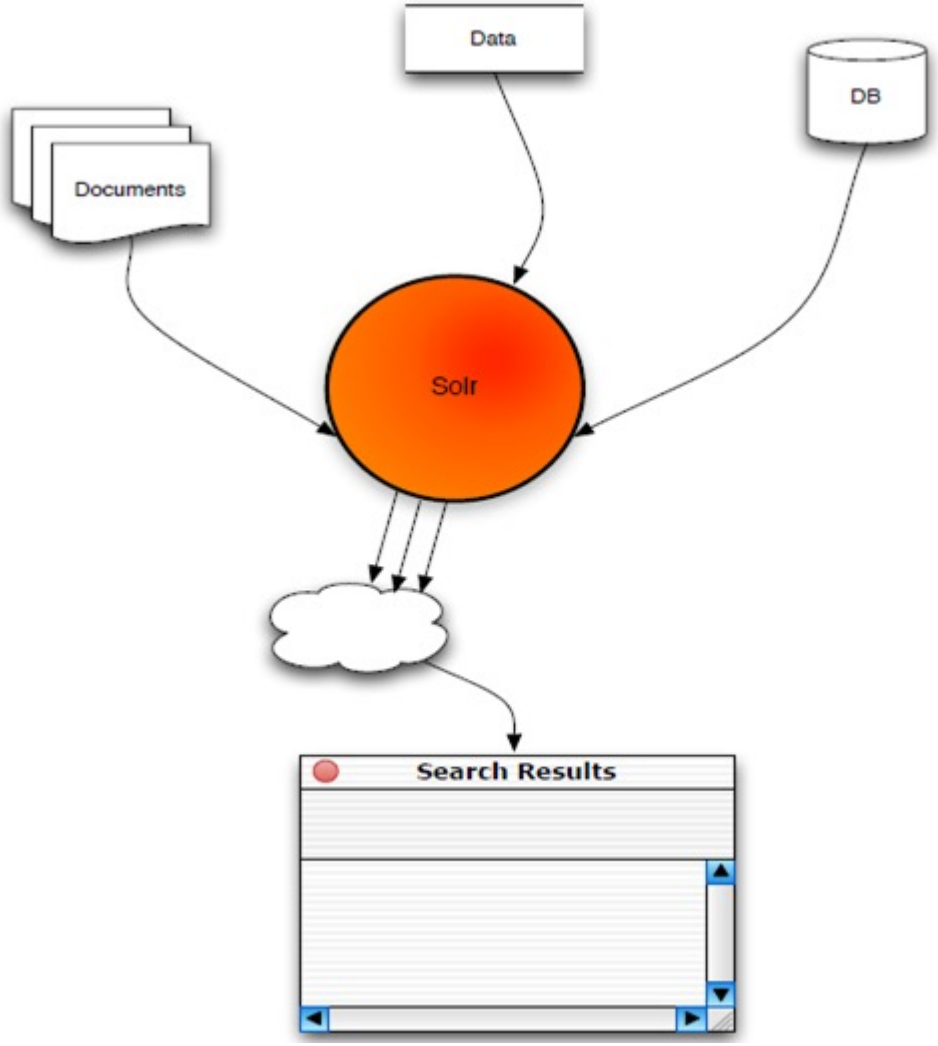

Source: http://www.slideshare.net/erikhatcher/solr-search-at-the-speed-of-light

Apache Solr - Features

- Full-text search
- Faceted search (similar to groupby clause in RDBMS)
- Scalability
  - Caching
  - Replication
  - Distributed search
- Near real-time indexing
- Geospatial search
- And many more : highlighting, database integration, rich document (e.G., Word, PDF) handling

Solr – schema.xml

- Types with index and query Analyzers - similar to data type
- Fields with name, type and options
- Unique Key : Unique Identifier of a document. For e.g. “id”
- Dynamic Fields : *Dynamic fields* allow Solr to index fields that you did not explicitly define in your schema. For e.g. fieldName: *_i or *_txts
- Copy Fields : Solr has a mechanism for making copies of fields so that you can apply several distinct field types to a single piece of incoming information. field ‘a‘ populates field ‘b’ with its value before tokenizing (having different analyzer/filter).

Solr – Content Analysis

- Field Attributes
  - Name : Name of the field
  - Type : Data-type (FieldType) of the field
  - Indexed : Should it be indexed (indexed="true/false")
  - Stored : Should it be stored (stored="true/false")
  - Required : is it a mandatory field (required="true/false")
  - Multi-Valued : Would it will contains multiple values e.g. text: pizza, food (multiValued="true/false")

e.g. <field name="id" type="string" indexed="true" stored="true" required="true" multiValued="false" />

Solr – solrconfig.xml

- Data dir: where all index data will be stored
- Index configuration
- Cache configurations
- Request Handler configuration
- Search components, response writers, query parsers

Query Types

- Single and multi term queries
  - ex fieldname:value or title: software engineer
- +, -, AND, OR NOT operators.
  - ex. title: (software AND engineer)
- Range queries on date or numeric fields,
  - ex: timestamp: [* TO NOW] or price: [1 TO 100]
- Boost queries:
  - e.g. title:Engineer ^1.5 OR text:Engineer
- Fuzzy search : is a search for words that are similar in spelling
  - e.g. roam~0.8 => noam
- Proximity Search : with a sloppy phrase query. The close together the two terms appear, higher the score.
  - ex “apache lucene”~20 : will look for all dcuments where “apache” word occurs within 20 words of “lucene”

Solr/Lucene Use-cases

- Search
- Analytics
- NoSQL datastore
- Auto-suggestion / Auto-correction
- Recommendation Engine (MoreLikeThis)
- Relevancy Engine (Feedback to other applications)
- Solr as a White-List
- GeoSpatial based Search

Search

- **Application**
  - Eclipse, Hibernate search
- **E-Commerce** :
  - Flipkart.com, Infibeam.com, Buy.com, Netflix.com, ebay.com
- **Jobs**
  - Indeed.com, Simplyhired.com, Naukri.com
- **Auto**
  - AOL.com
- **Travel**
  - Cleartrip.com
- **Social Network**
  - Twitter.com, LinkedIn.com, mylife.com

Source: http://www.quora.com/Which-major-companies-are-using-Solr-for-search

Search (Contd.)

- **Search Engine**
  - Yandex.ru, DuckDuckGo.com
- **News Paper**
  - Guardian.co.uk
- **Music/Movies**
  - Apple.com, Netflix.com
- **Events**
  - Stubhub.com, Eventbrite.com
- **Cloud Log Management**
  - Loggly.com
- **Others**
  - Whitehouse.gov

Faceting

- Grouping results based on field value
- Facet on: field terms, queries, date ranges
- &facet=on &facet.field=job_title &facet.query=salary:[30000 TO 100000]
- http://wiki.apache.org/solr/SimpleFacetParameters

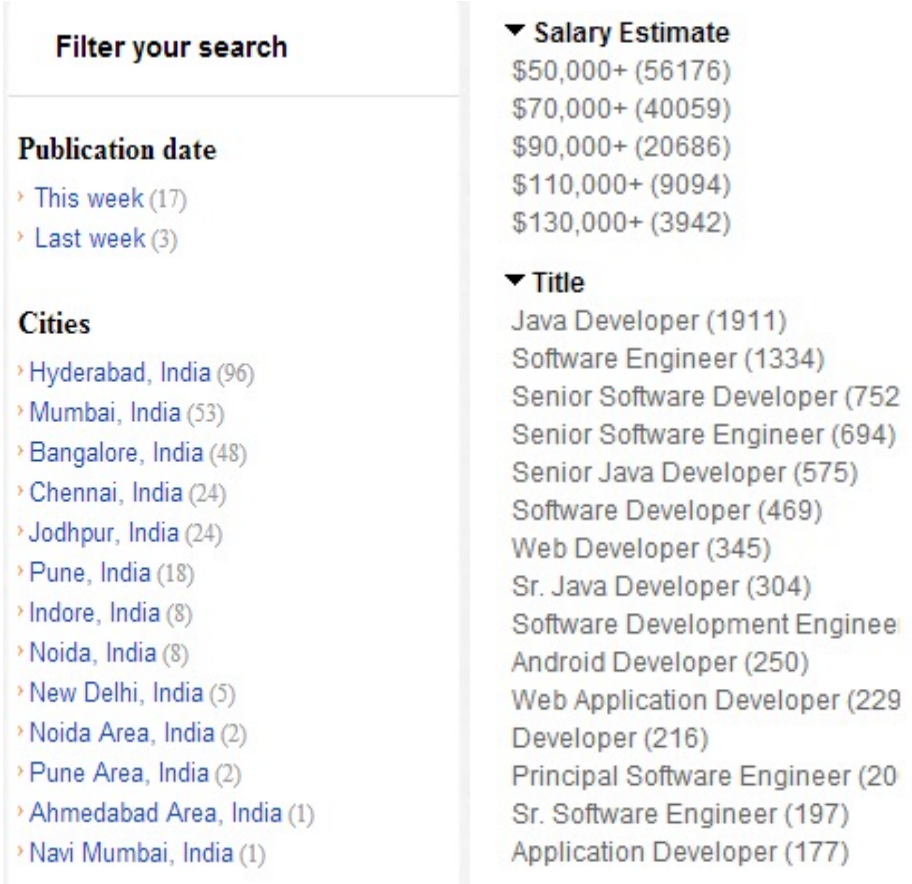

Source: www.career9.com, www.indeed.com

Analytics

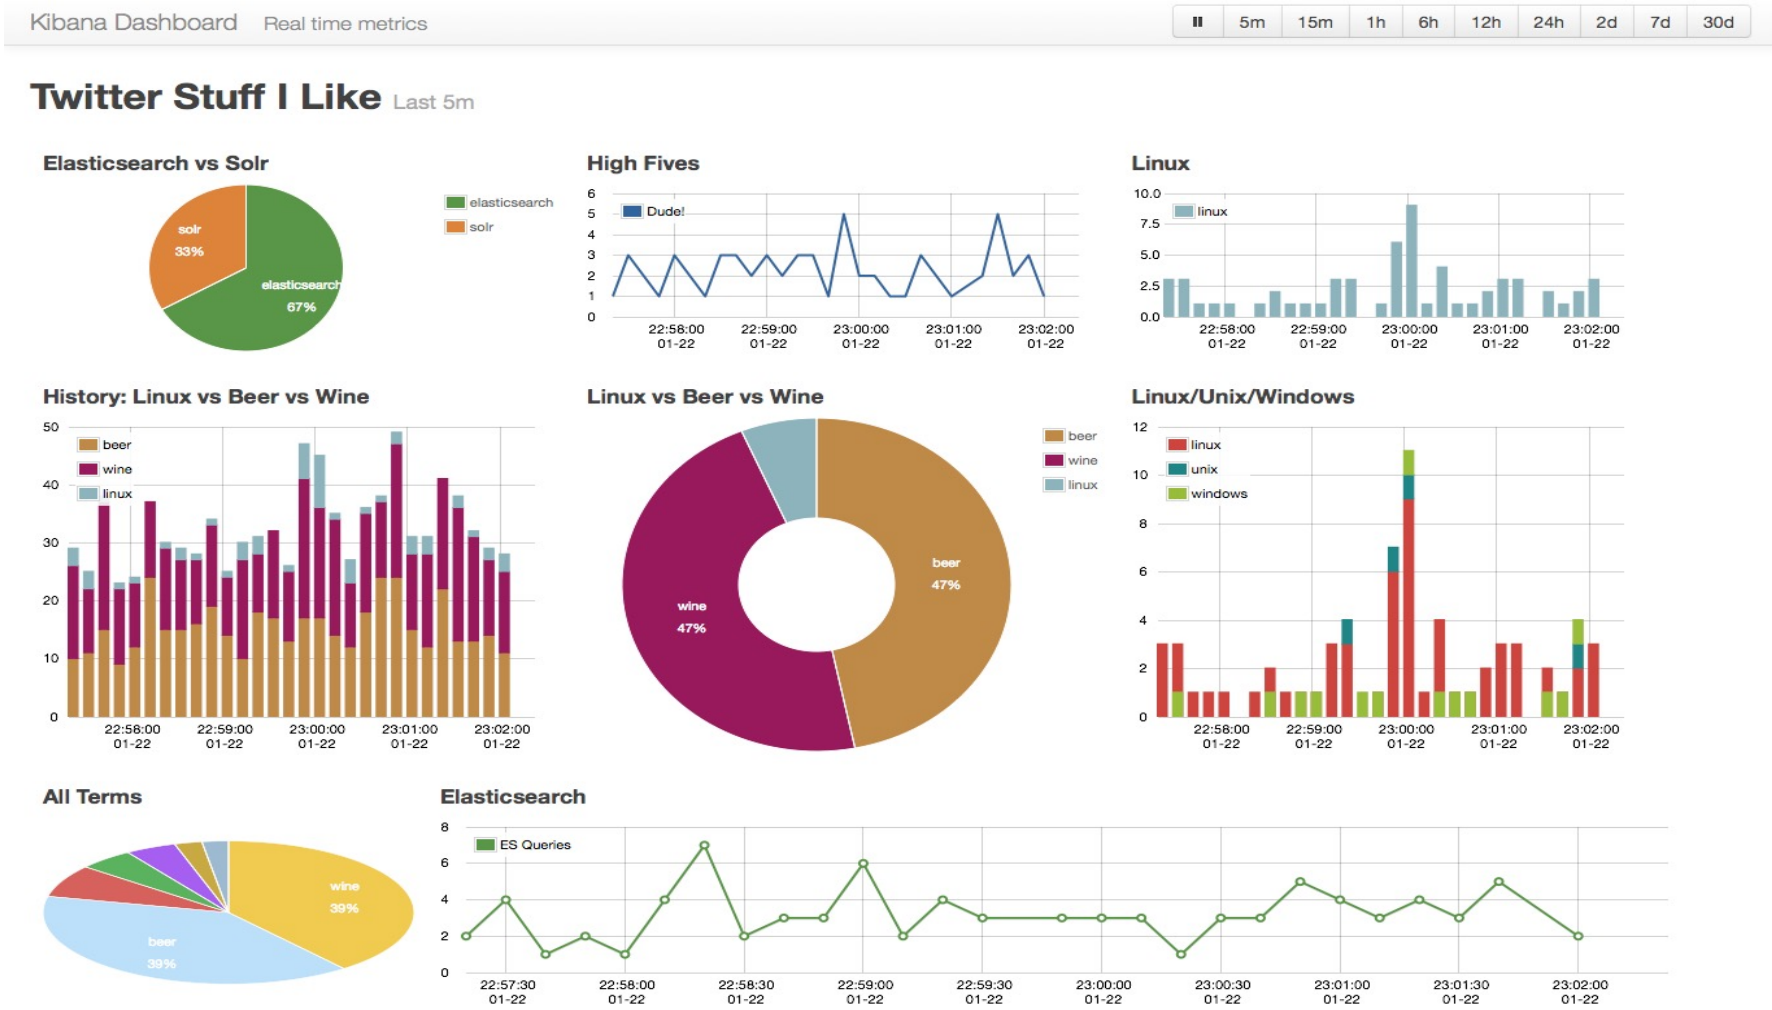

- Analytics source : Kibana.org based on ElasticSearch and Logstash
- Image Source : http://semicomplete.com/presentations/logstash-monitorama-2013/#/8

Autosuggestion

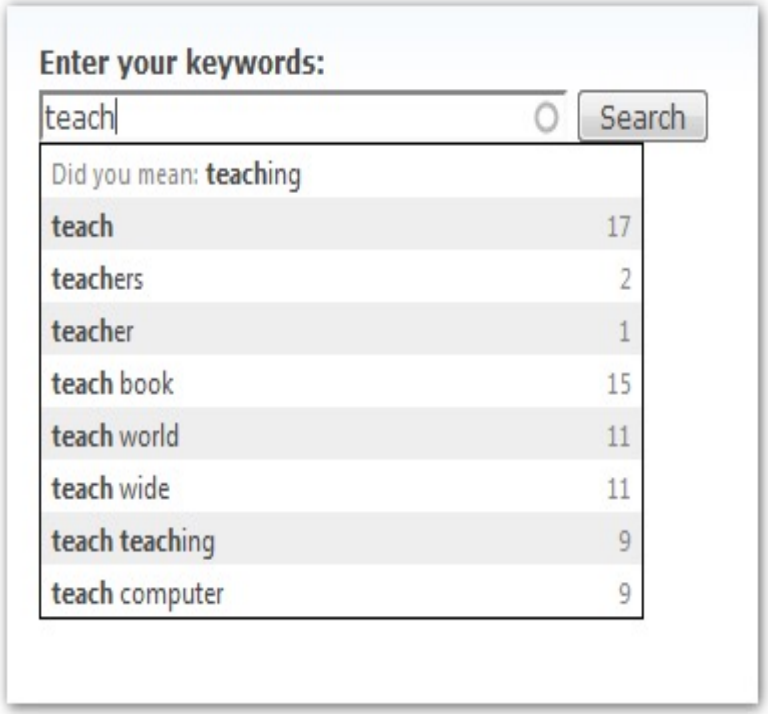

Source: www.drupal.org , www.yelp.com

Integration

- Clustering (Solr-Carrot2)
- Named Entity extraction (Solr-UIMA)
- SolrCloud (Solr-Zookeeper)
- Parsing of many Different File Formats (Solr-Tika)
- Machine Learning/Data Mining (Apache Mahout)
- Large scale Indexing (Hadoop)

SolrCtl Command

- The solrctl utility is a wrapper shell script included with Cloudera Search for managing collections, instance directories, configs, Apache Sentry permissions, and more.

Syntax

The general `solrctl` command syntax is:

```
solrctl [options] command [command-arg] [command [command-arg]] ...
```

Source: https://www.cloudera.com/documentation/enterprise/5-14-x/topics/search_solrctl_ref.html

SolrCtl Collection Commands

```
collection [--create <name> -s <numShards>
                           [-a]
                           [-c <configName>]
                           [-r <replicationFactor>]
                           [-m <maxShardsPerHost>]
                           [-n <createHostSet>]]
           [--delete <name>]
           [--reload <name>]
           [--stat <name>]
           [--deletedocs <name>]
           [--list]
           [--create-snapshot <snapshotName> -c <collectionName>]
           [--delete-snapshot <snapshotName> -c <collectionName>]
           [--list-snapshots <collectionName>]
           [--describe-snapshot <snapshotName> -c <collectionName>]
           [--prepare-snapshot-export <snapshotName> -c <collectionName> -d <destDir> [
           [--export-snapshot <snapshotName> [-s <sourceDir>] [-c <collectionName>] -d
           [--restore name  -b <backupName> -l <backupLocation> -i <requestId>
                            [-a]
                            [-c <configName>]
                            [-r <replicationFactor>]
                            [-m <maxShardsPerNode>]]
           [--request-status <requestId>]
```

Source: https://www.cloudera.com/documentation/enterprise/5-14-x/topics/search_solrctl_ref.html

- solrctl collection –list

Lists the collection

- solrctl config --create logs_config predefinedTemplate -p immutable=false

logs_config => config name

predefinedTemplate => existing config template

- solrctl instancedir --generate $HOME/logs_config
- solrctl collection --create logNew2 -c logs_config

logNew2 => collection Name

Update or add data to collection

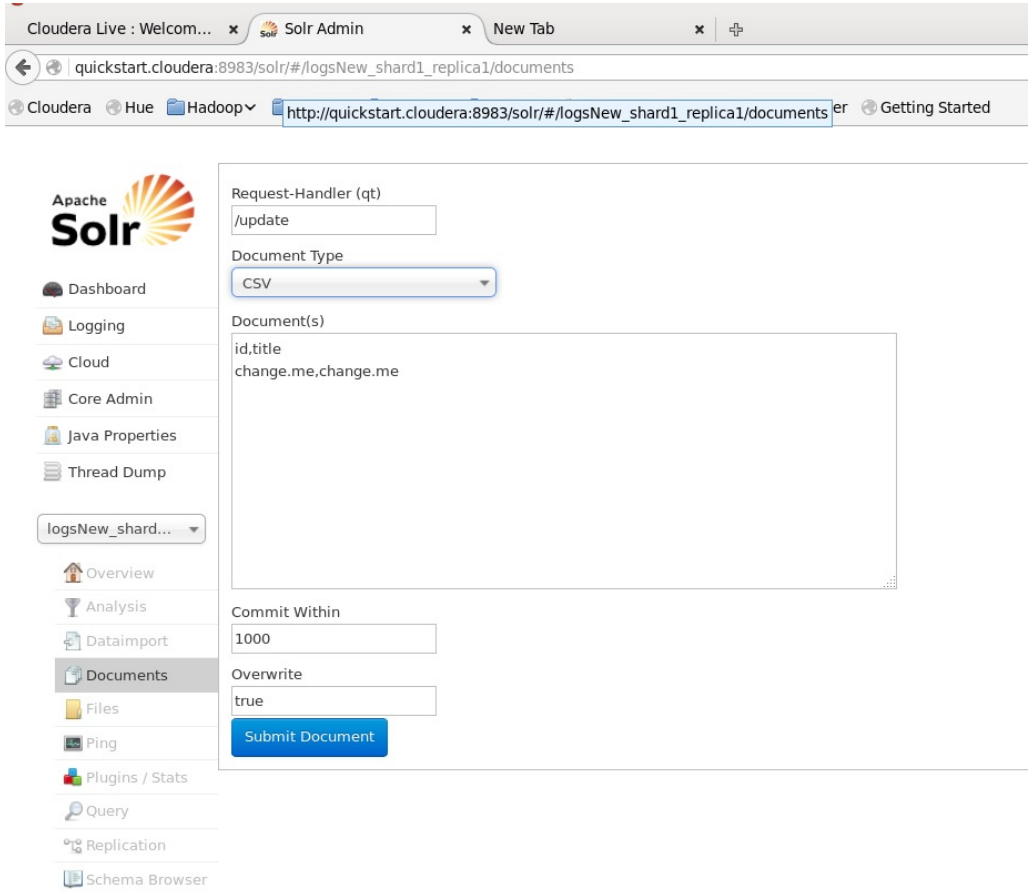

Query Syntax

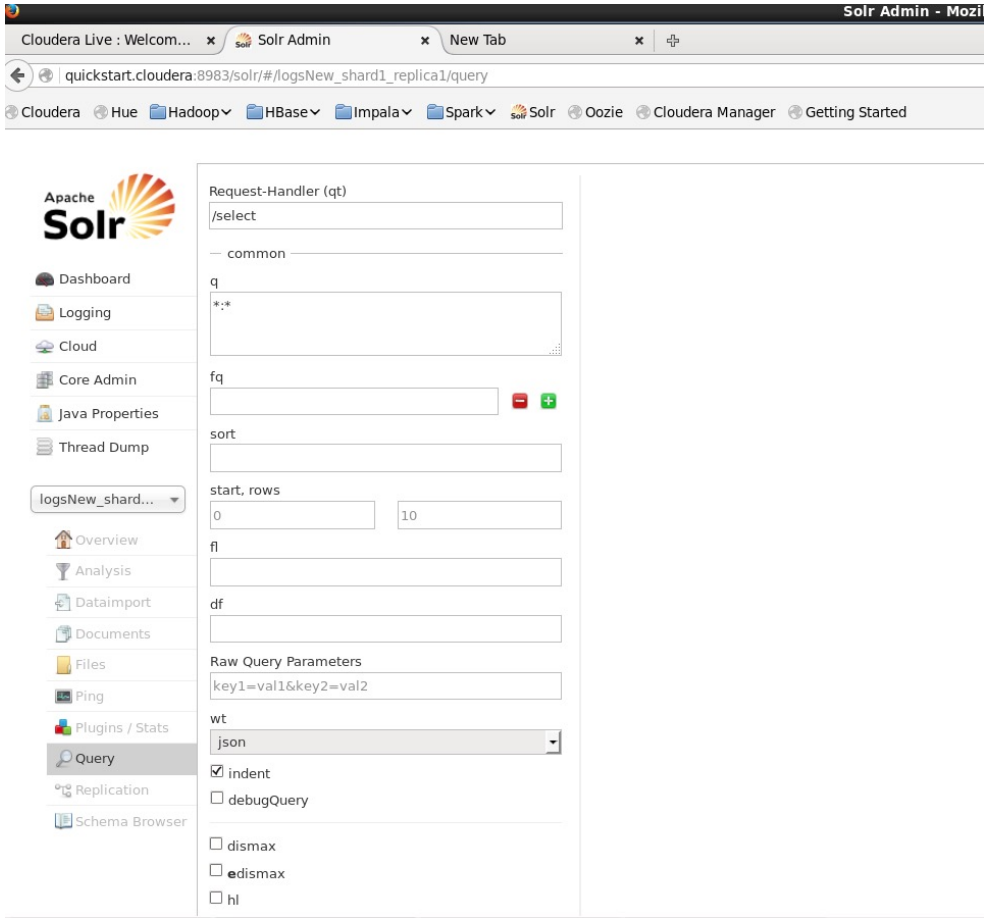

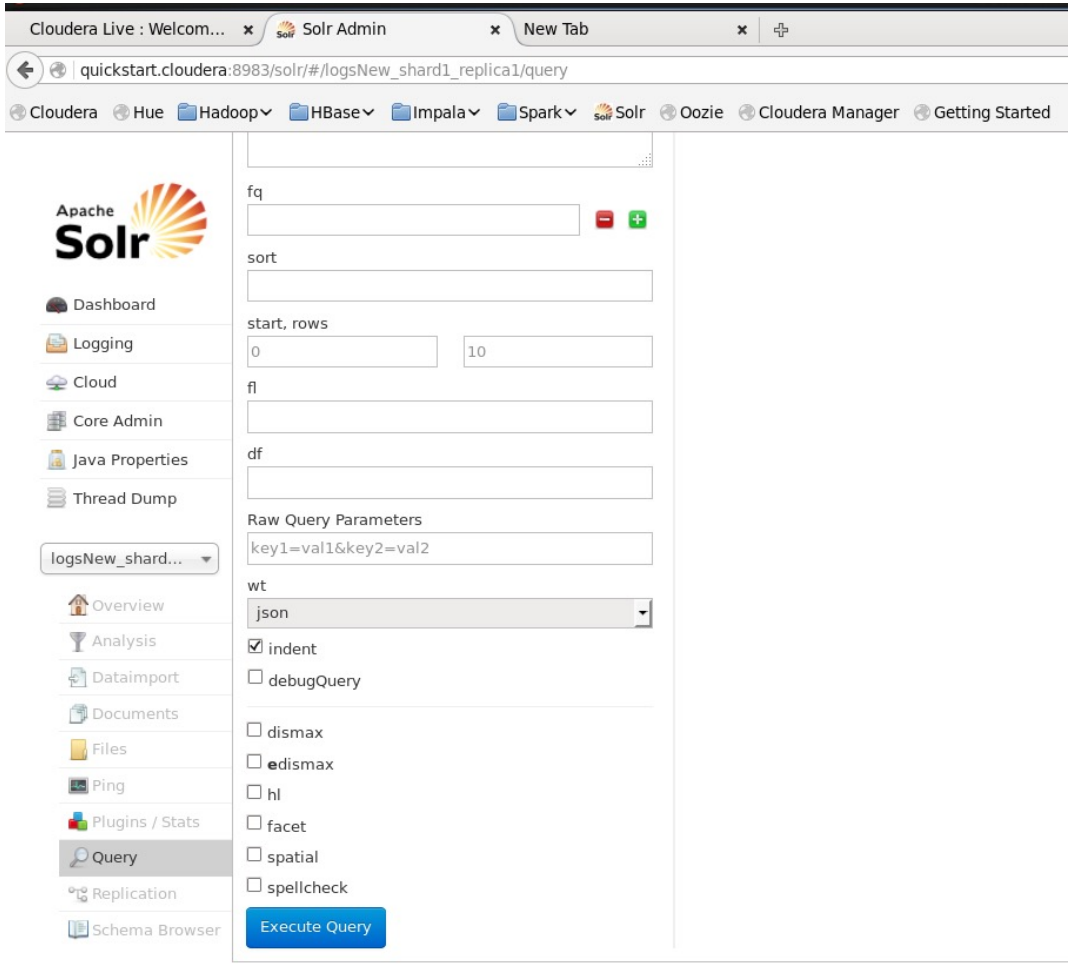

Results

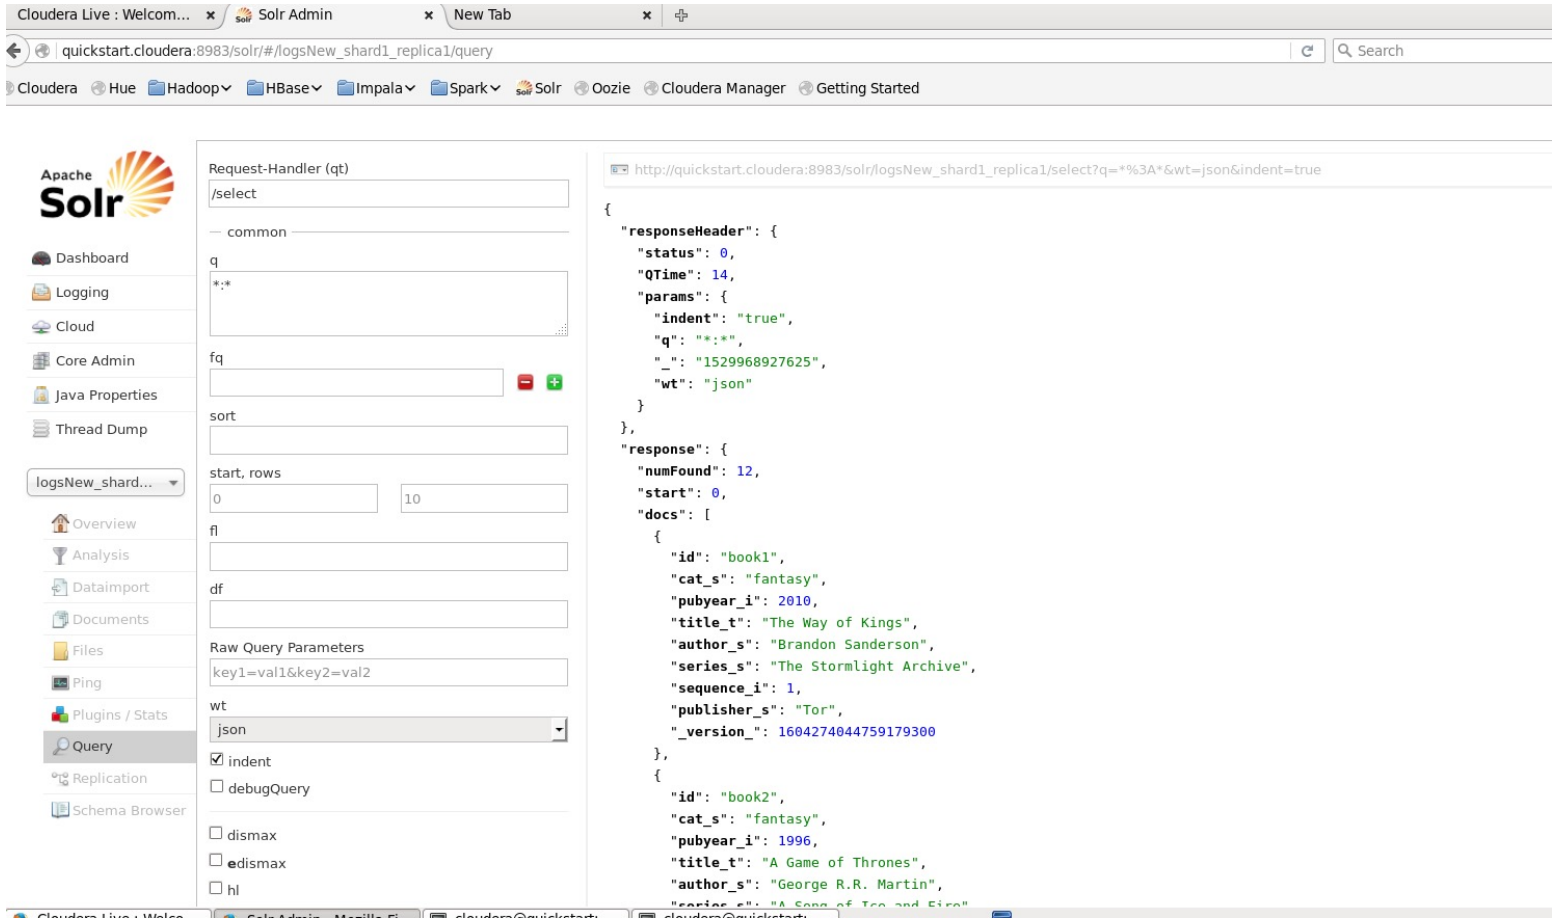
Cloudera Live : Welcom...
Solr Admin
New Tab
quickstart.cloudera:8983/solr/#/logsNew_shard1_replica1/query
Search
Cloudera
Hue
Hadoop
HBase
Impala
Spark
Solr
Oozie
Cloudera Manager
Getting Started
Apache Solr
Dashboard
Logging
Cloud
Core Admin
Java Properties
Thread Dump
logsNew_shard...
Overview
Analysis
Dataimport
Documents
Files
Ping
Plugins / Stats
Query
Replication
Schema Browser
Request-Handler (qt)
/select
common
q
:
fq
sort
start, rows
0
10
fl
df
Raw Query Parameters
key1=val1&key2=val2
wt
json
indent
debugQuery
dismax
edismax
hl
http://quickstart.cloudera:8983/solr/logsNew_shard1_replica1/select?q=*%3A*&wt=json&indent=true
{
"responseHeader": {
"status": 0,
"QTime": 14,
"params": {
"indent": "true",
"q": "*:*",
"_": "1529968927625",
"wt": "json"
}
},
"response": {
"numFound": 12,
"start": 0,
"docs": [
{
"id": "book1",
"cat_s": "fantasy",
"pubyear_i": 2010,
"title_t": "The Way of Kings",
"author_s": "Brandon Sanderson",
"series_s": "The Stormlight Archive",
"sequence_i": 1,
"publisher_s": "Tor",
"_version_": 1604274044759179300
},
{
"id": "book2",
"cat_s": "fantasy",
"pubyear_i": 1996,
"title_t": "A Game of Thrones",
"author_s": "George R.R. Martin",

Creating Schema Config

- solrctl instancedir --generate /tmp/films

Editing Schema Config

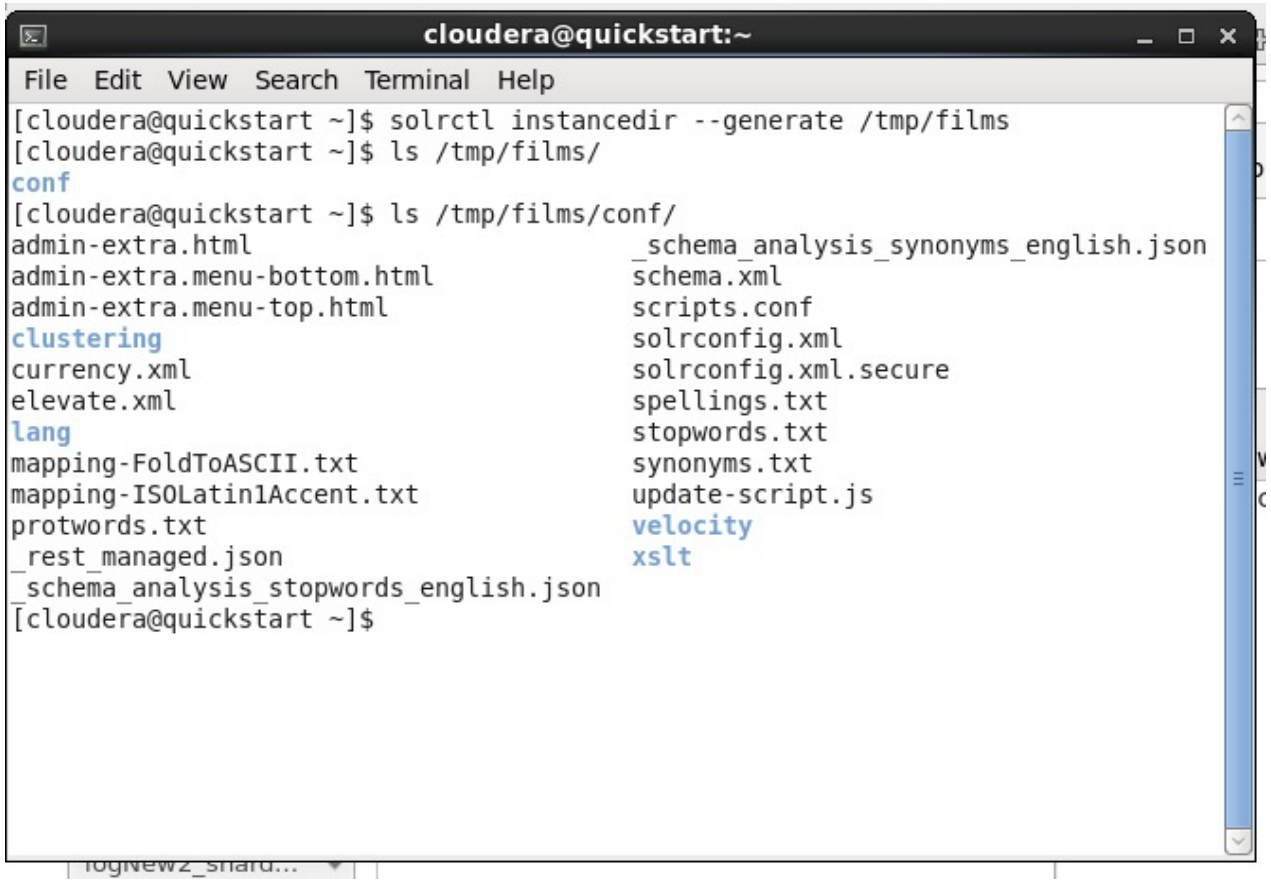

Editing Schema Config :
gedit /tmp/films/conf/schema.xml

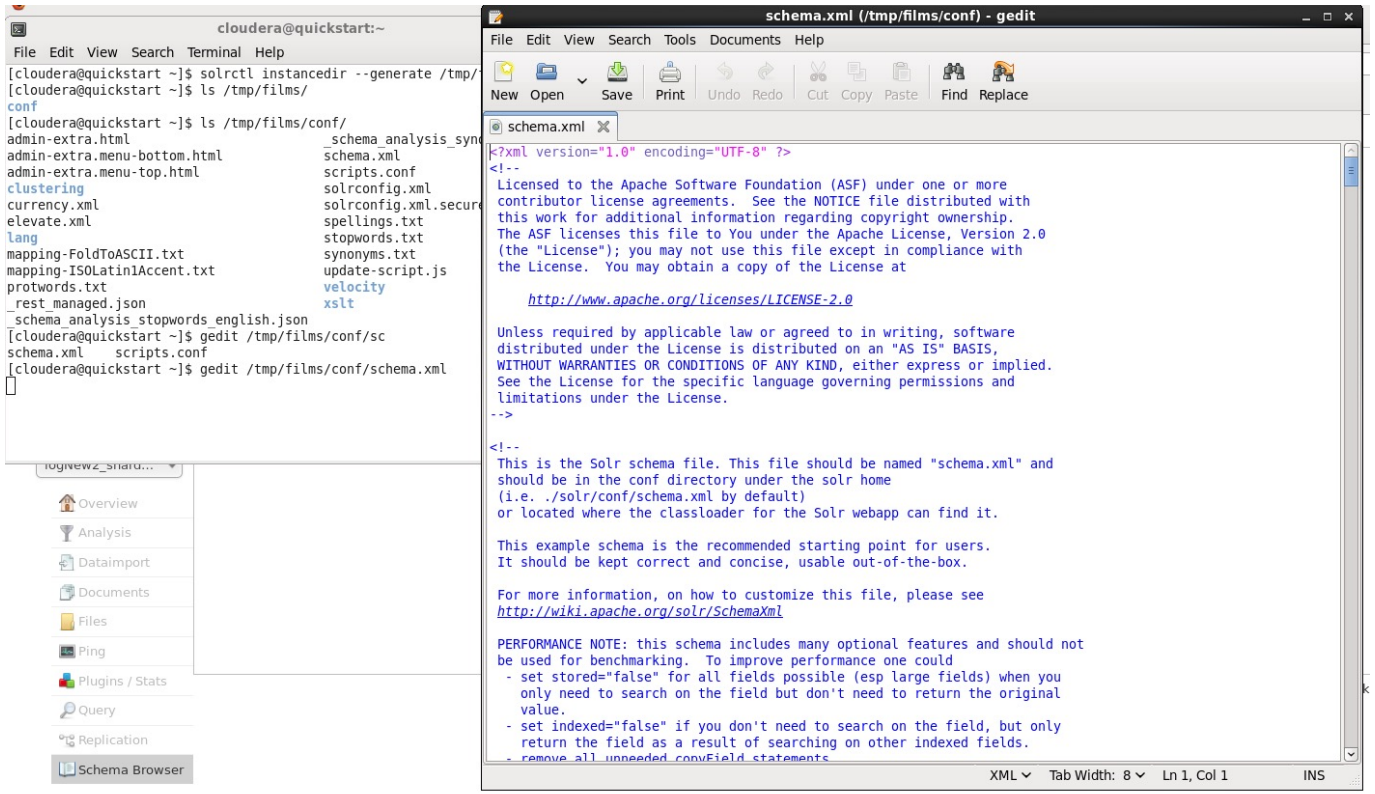

Creating new fieldType (Use this to create your own directed_by field)

- http://www.solrtutorial.com/schema-xml.html
- <fields>
  <field name="id" type="string" indexed="true" stored="true" required="true" />
  <field name="name" type="textgen" indexed="true" stored="true"/>
  ...
  </fields>

Film Dataset commands

Instancedir and collection

- solrctl instancedir --create films /tmp/films
- solrctl collection --create films

Edit Schema

- ls /tmp/films/conf/
- gedit /tmp/films/conf/schema.xml

References

- http://www.lucenetutorial.com/lucene-vs-solr.html
- https://lucene.apache.org/solr/
- https://lucene.apache.org/solr/guide/6_6/the-standard-query-parser.html
- https://lucene.apache.org/solr/guide/8_5/solr-tutorial.html

Commands Details

- solrctl instancedir -- generate - Use this command to generate new instance.
- solrctl instancedir -- create <collection_name> - To upload the contents of instance directory to Zookeeper.
- solrctl collection -- create <collection_name> - Used to create new collection.

Commands

- solrctl config --create logs_config predefinedTemplate -p immutable=false
- solrctl instancedir --generate $HOME/logs_config
- solrctl collection --create logNew2 -c logs_config
- solrctl instancedir --generate /tmp/films
- ls /tmp/films/conf