CSCE 5300

Spark Streaming

What is Streaming?

- Data Streaming is a technique for transferring data so that it can be processed as a steady and continuous stream

- Streaming technologies are becoming increasingly important with the growth of the Internet

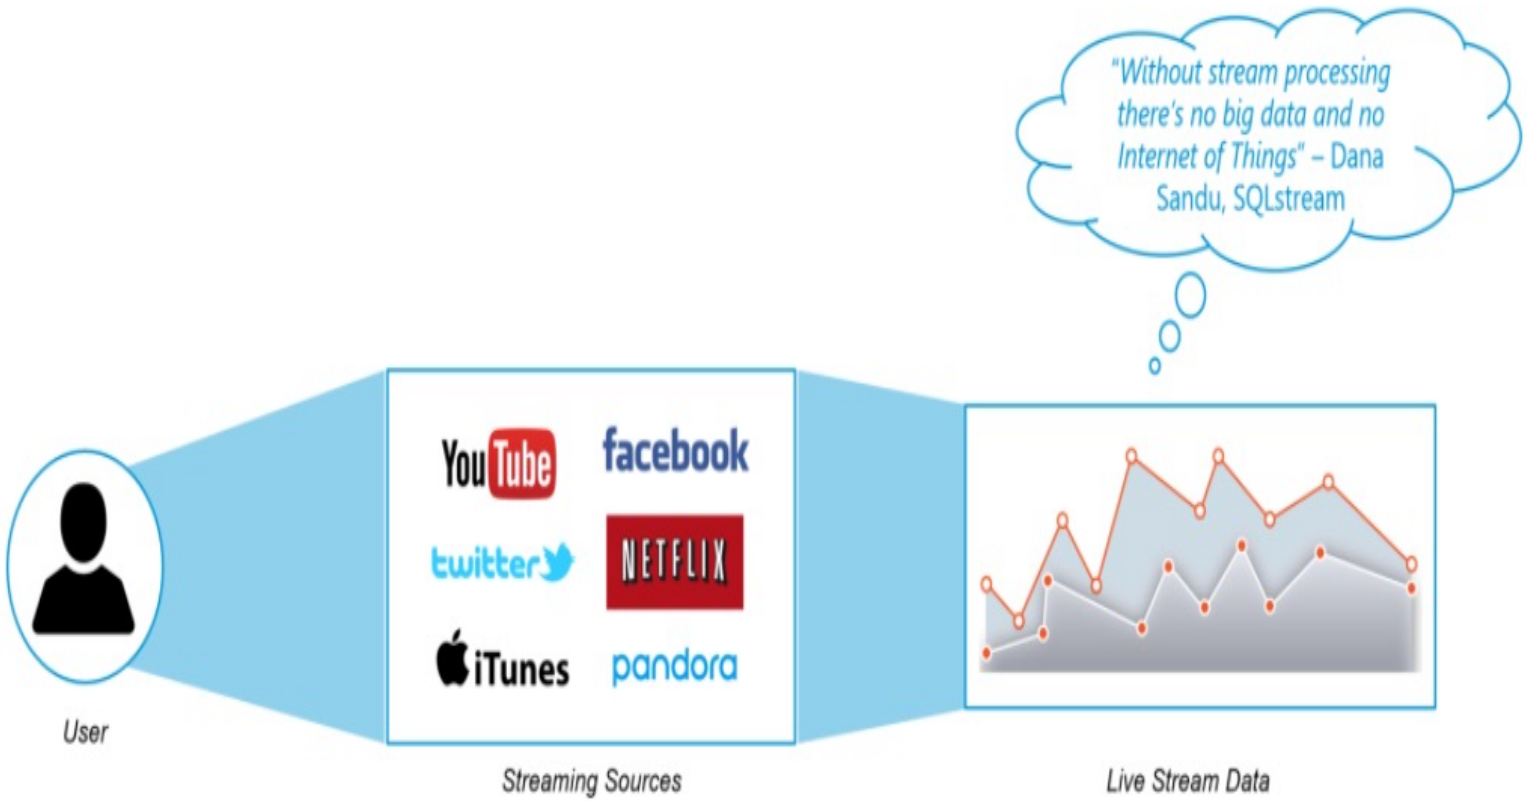

Figure: *What is Streaming?*

Why Spark Streaming?

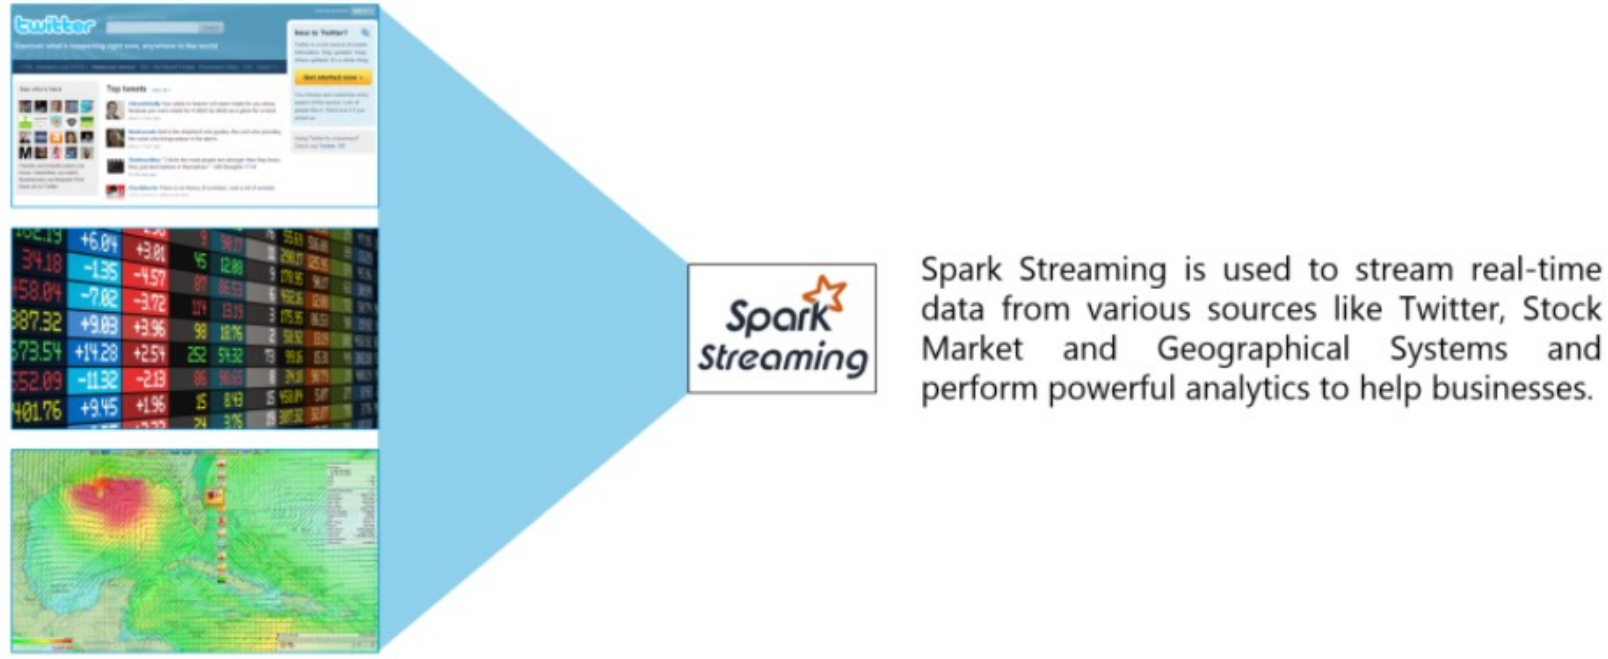

Spark Streaming is used to stream real-time data from various sources like Twitter, Stock Market and Geographical Systems and perform powerful analytics to help businesses.

Figure: *Why Spark Streaming?*

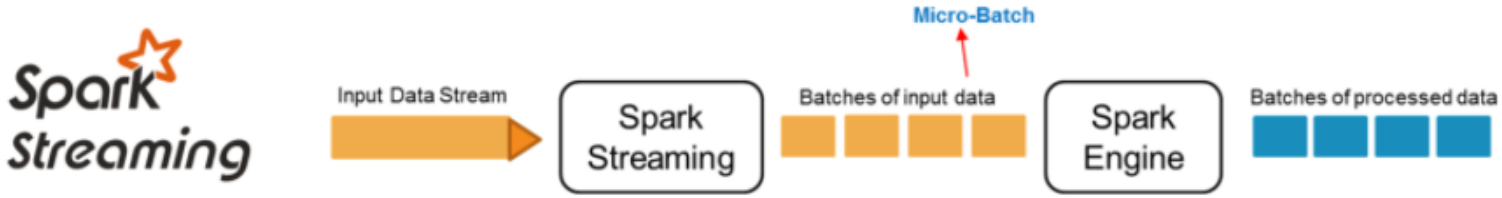

Figure: *Streams in Spark Streaming*

Spark Streaming Features

- **Scaling:** Spark Streaming can easily scale to hundreds of nodes.
- **Speed:** It achieves low latency.
- **Fault Tolerance:** Spark has the ability to efficiently recover from failures.
- **Integration:** Spark integrates with batch and real-time processing.
- **Business Analysis:** Spark Streaming is used to track the behavior of customers which can be used in business analysis

Spark Streaming Workflow

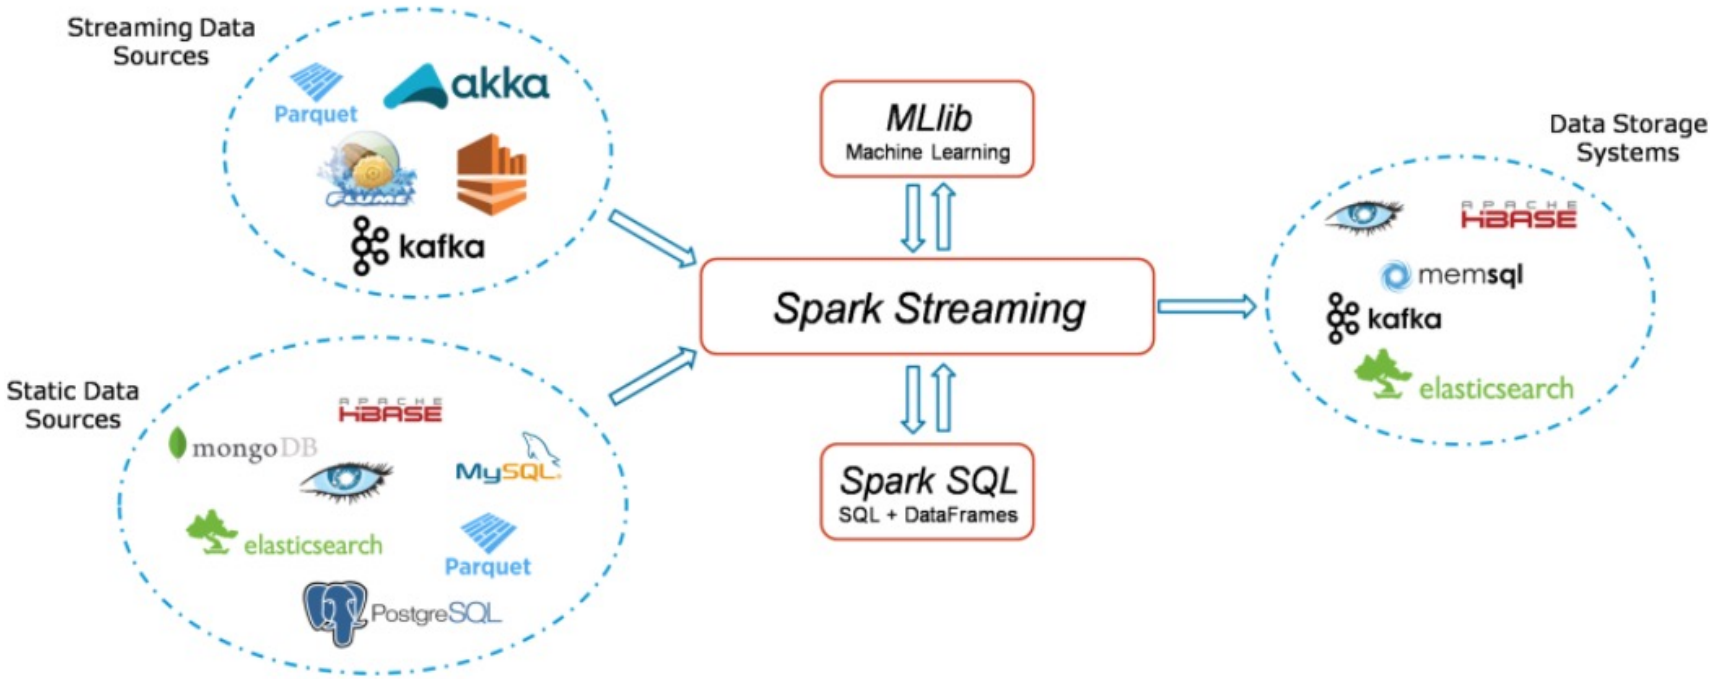

Figure: *Overview Of Spark Streaming*

Spark Streaming Fundamentals

- Streaming Context
- DStream
- Caching
- Accumulators, Broadcast Variables and Checkpoints

Streaming Context

- *Streaming Context* consumes a stream of data in Spark.
- It registers an *Input DStream* to produce a *Receiver* object.
- It is the main entry point for Spark functionality.
- Spark provides a number of default implementations of sources like Twitter, Akka Actor and ZeroMQ that are accessible from the context.

Figure: *Spark Streaming Context*

Figure: *Default Implementation Sources*

DStream

- *Discretized Stream* (DStream) is the basic abstraction provided by Spark Streaming.
- It is a continuous stream of data.
- It is received from a data source or a processed data stream generated by transforming the input stream.

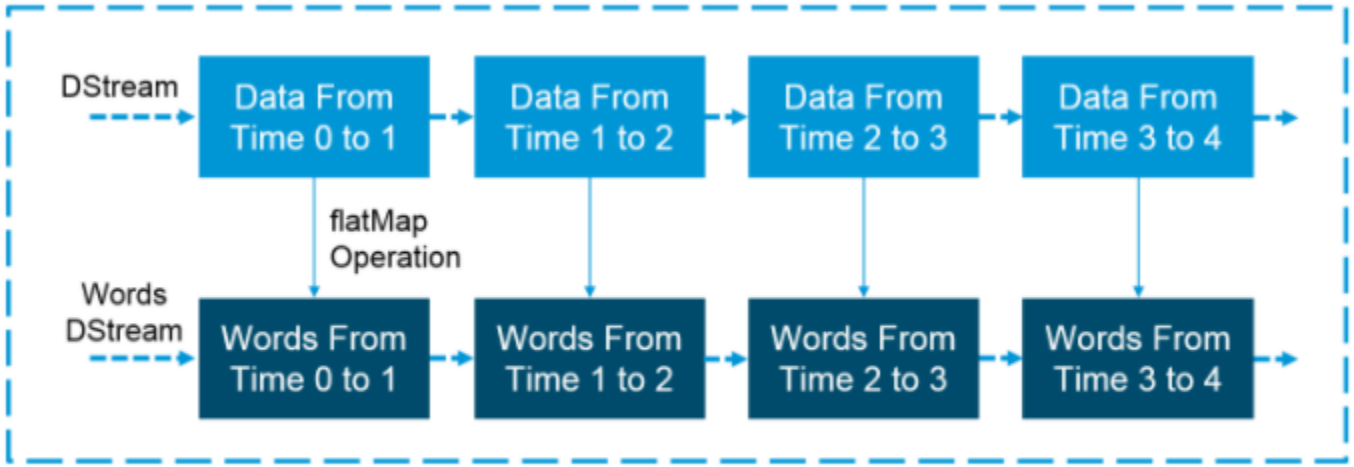

Figure: *Extracting words from an Input DStream*

Transformations on DStreams

map(*func*)	map(*func*) returns a new DStream by passing each element of the source DStream through a function *func*.
flatMap(*func*)	flatMap(*func*) is similar to map(*func*) but each input item can be mapped to 0 or more output items and returns a new DStream by passing each source element through a function *func*.
filter(*func*)	filter(*func*) returns a new DStream by selecting only the records of the source DStream on which *func* returns true.
reduce(*func*)	reduce(*func*) returns a new DStream of single-element RDDs by aggregating the elements in each RDD of the source DStream using a function *func*.
groupBy(*func*)	groupBy(*func*) returns the new RDD which basically is made up with a key and corresponding list of items of that group.

Accumulators, Broadcast Variables and Checkpoints

- **Accumulators:**

➢ *Accumulators* are variables that are only added through an associative and commutative operation.

➢ They are used to implement counters or sums. Tracking accumulators in the UI can be useful for understanding the progress of running stages.

➢ Spark natively supports numeric accumulators. We can create named or unnamed accumulators.

- **Broadcast Variables:**

➢ *Broadcast variables* allow the programmer to keep a read-only variable cached on each machine rather than shipping a copy of it with tasks.

➢ They can be used to give every node a copy of a large input dataset in an efficient manner.

➢ Spark also attempts to distribute broadcast variables using efficient broadcast algorithms to reduce communication cost.

- **Checkpoints:** *Checkpoints* are similar to checkpoints in gaming. They make it run 24/7 and make it resilient to failures unrelated to the application logic.

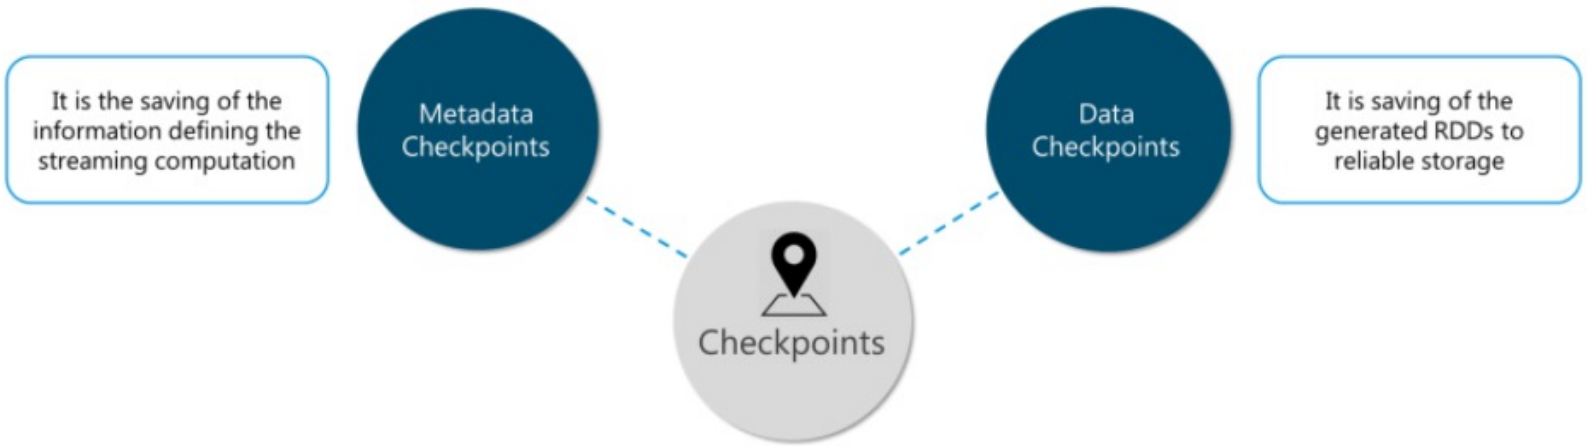

Figure: *Features of Checkpoints*

Spark Streaming

File Generation

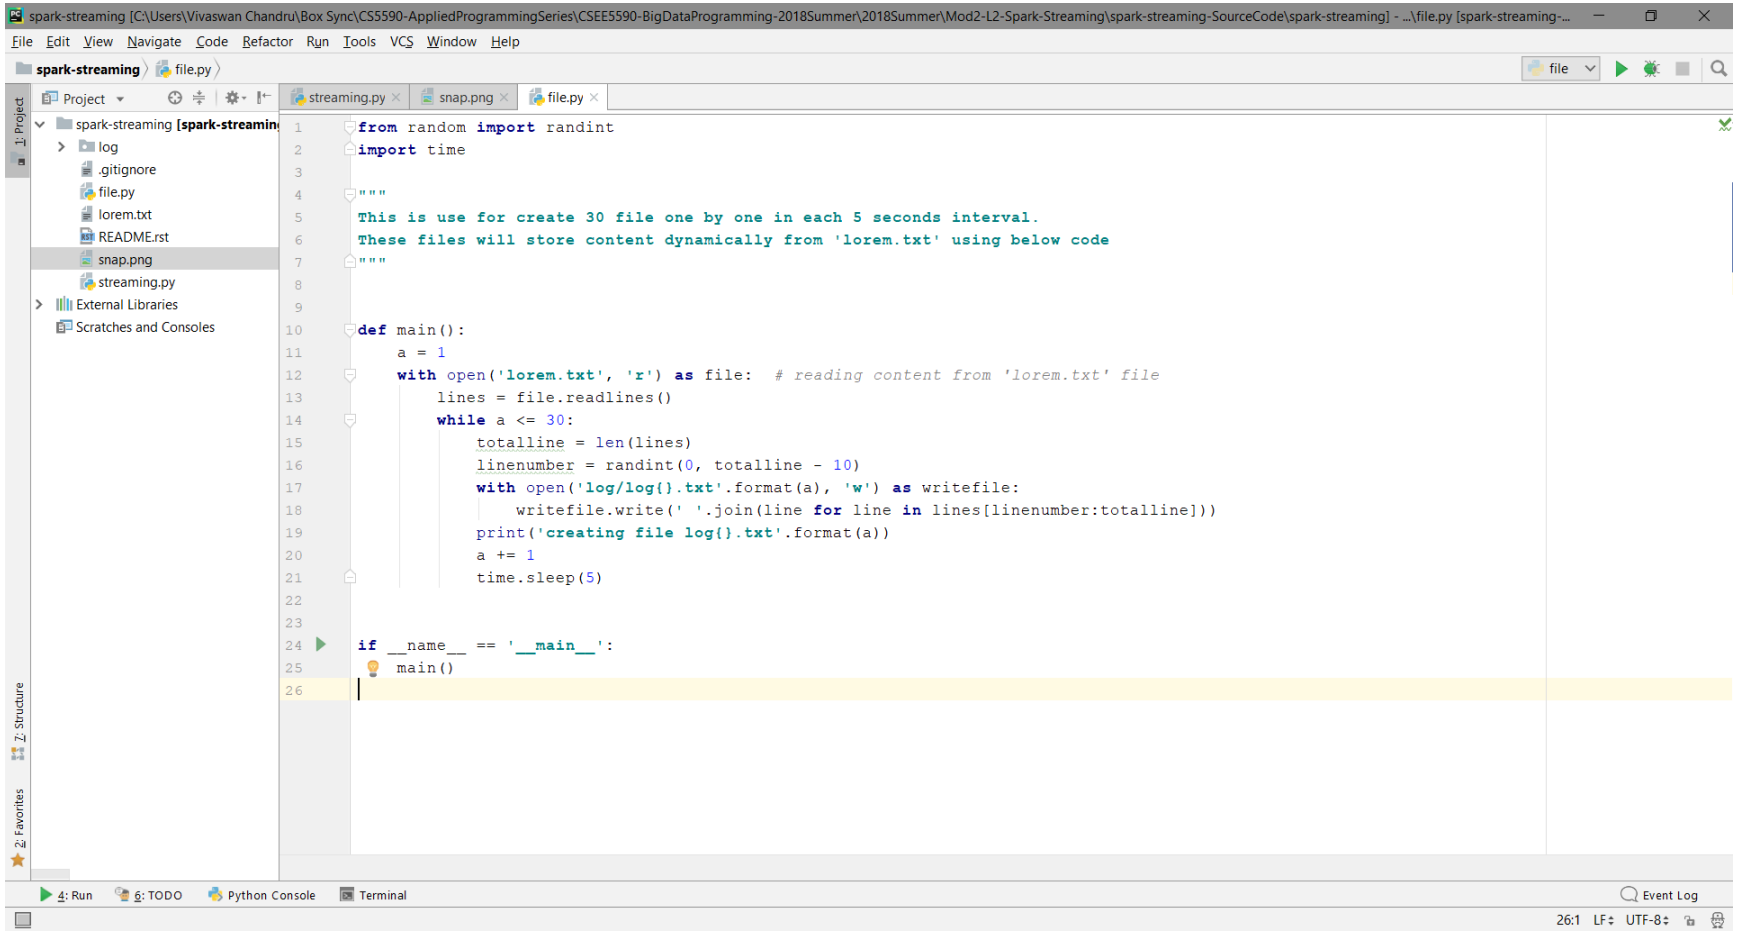

Download Netcat

- Go https://joncraton.org/files/nc111nt.zip
- Unzip the files and use nc as password
- Go to environment variables
- Select Path and Add new path as you location to netcat folder.
- Open command line
- Type nc to check either netcat is working or not

Running Results

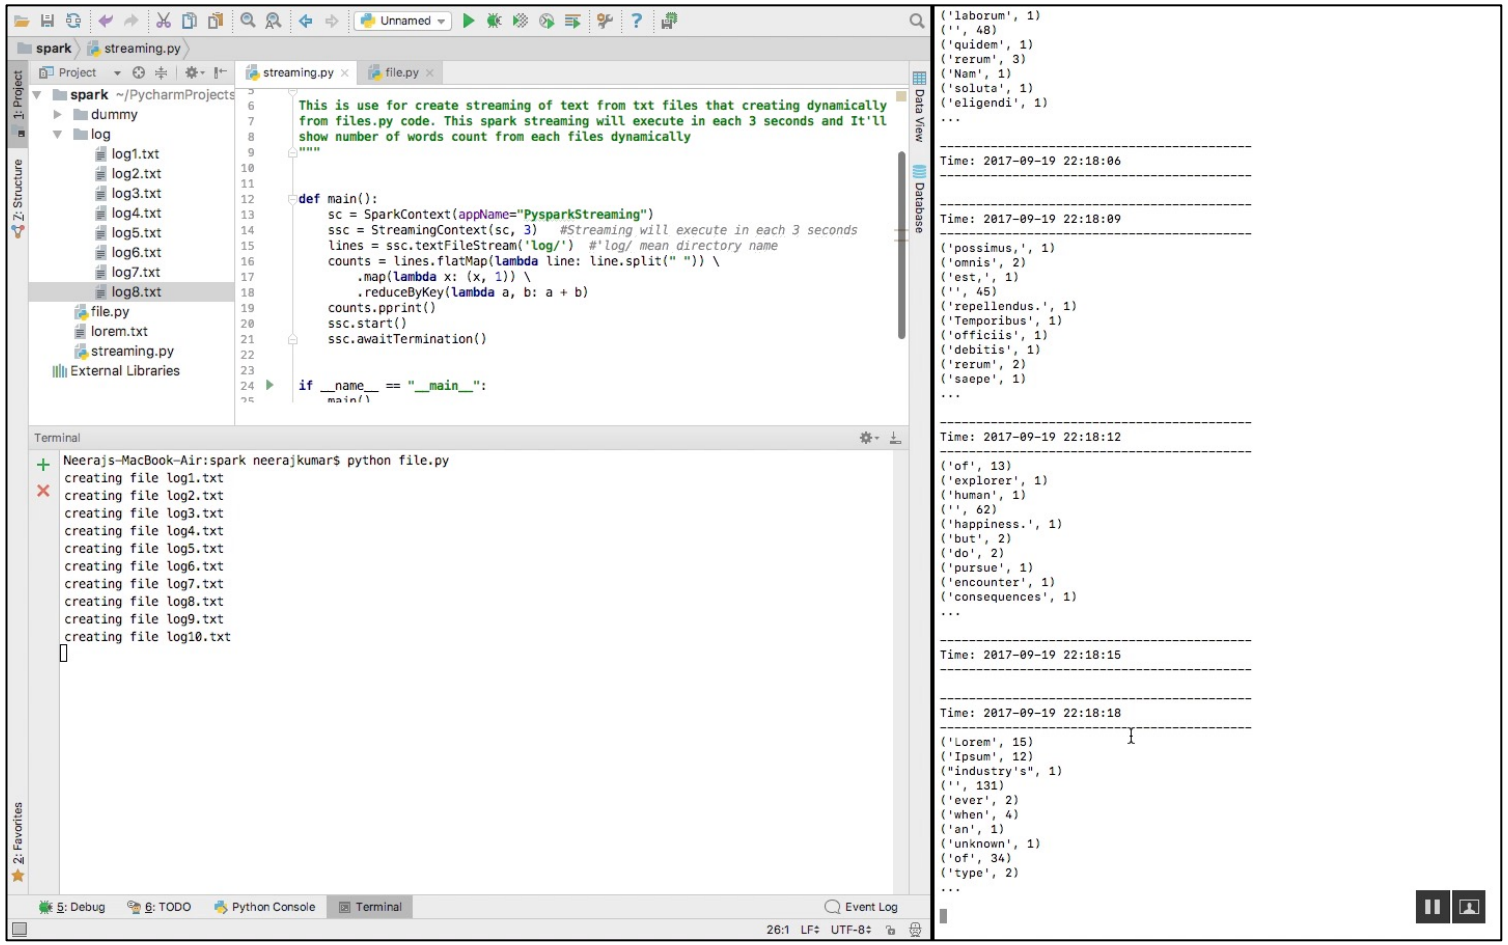

References

- https://spark.apache.org/docs/2.2.0/streaming-programming-guide.html
- https://www.edureka.co/blog/spark-streaming/