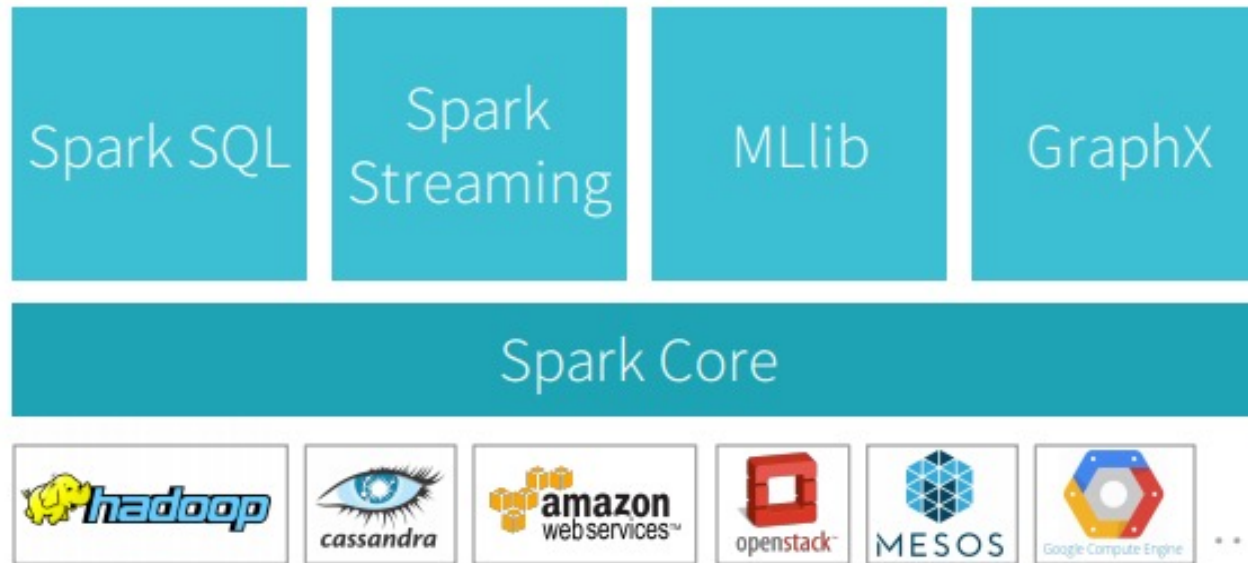


CSCE 5300

GraphX and GraphFrames

Apache Spark Engine

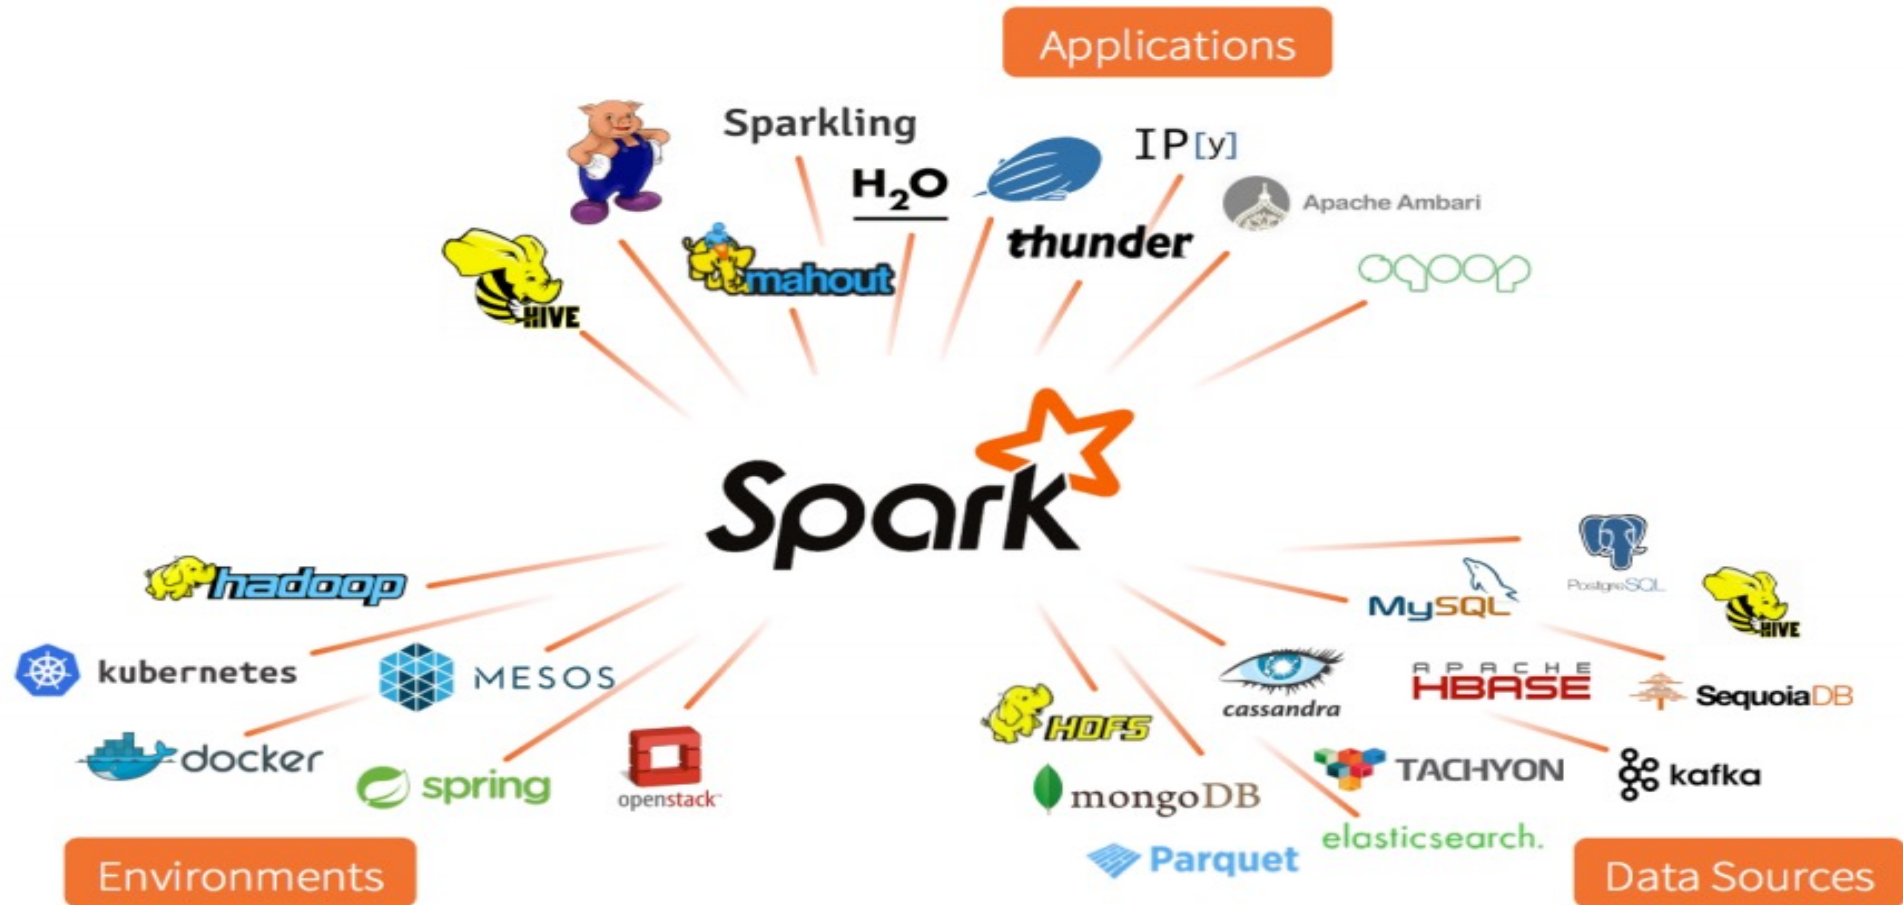


Scale out, fault tolerant

Python, Java, Scala, and R APIs

Standard libraries

Open source Eco-System




Graphs

- *GraphX is to RDDs as GraphFrames are to DataFrames*
- GraphFrames represent graphs: vertices (e.g., users) and edges (e.g., relationships between users).
- GraphFrames are based upon Spark DataFrames
- GraphX are based upon RDDs

GraphX vs GraphFrames

	GraphFrames	GraphX
Built on	DataFrames	RDDs
Languages	Scala, Java, Python	Scala
Use cases	Queries & algorithms	Algorithms
Vertex IDs	Any type (in Catalyst)	Long
Vertex/edge attributes	Any number of DataFrame columns	Any type (VD, ED)
Return types	GraphFrame or DataFrame	Graph[VD, ED], or RDD[Long, VD]

Continued ..



GraphFrames Configuration

- Download Spark (Latest Version)

<https://spark.apache.org/downloads.html>

From Command Line:

- Go to bin
 - Spark-shell –packages graphframes:graphframes:0.5.0-spark2.3-s_2.11

Add jars

- From project structure add the downloaded jars in intelliJ

Load Pyspark GraphFrames

From PyCharm:

```
import os  
os.environ["PYSPARK_SUBMIT_ARGS"] = (  
"--packages graphframes:graphframes:0.5.0-spark2.0-s_2.11 pyspark-shell")
```

From Command Line:

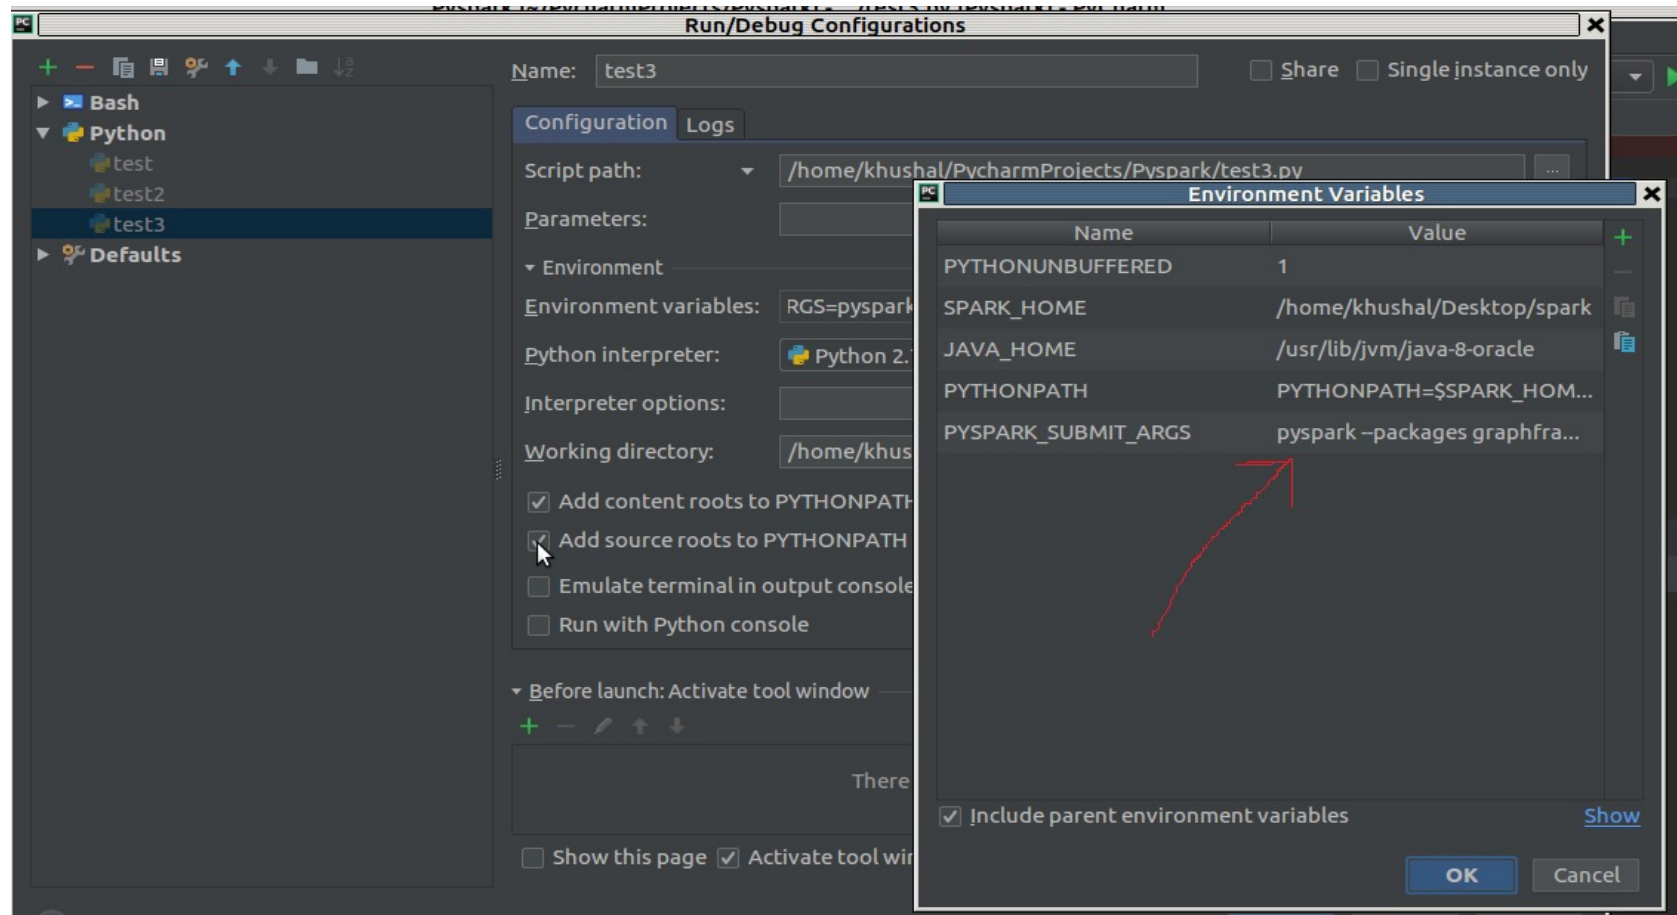
```
Pyspark --packages graphframes:graphframes:0.5.0-spark2.0-s_2.11 pyspark-shell
```


Load in PyCharm or Edit Configuration Settings

```
import os
os.environ["PYSPARK_SUBMIT_ARGS"] = ("pyspark --packages graphframes:graphframes:0.5.0-spark2.0-s_2.11")

from graphframes import *
from pyspark.sql import SparkSession
```

Continued ..



Pyspark command line

```
(venv) khushal@UDIC-GPU:~/PycharmProjects/Pyspark$ pyspark --packages graphframes:graphframes:0.5.0-spark2.0-s_2.11
```

Pyspark Loaded with graphframes

```
graphframes#graphframes;0.5.0-spark2.0-s_2.11 from spark-packages in [default]
org.scala-lang#scala-reflect;2.11.0 from central in [default]
org.slf4j#slf4j-api;1.7.7 from central in [default]
-----
|          conf          | modules          || artifacts |
|          number| search|dwnlded|evicted|| number|dwnlded|
-----
|          default      | 5 | 0 | 0 | 0 || 5 | 0 |
-----

:: retrieving :: org.apache.spark#spark-submit-parent-343fbdcf-5c2c-47d1-897d-b7f03882c087
  confs: [default]
  0 artifacts copied, 5 already retrieved (0kB/10ms)
2018-07-10 22:45:24 WARN  Utils:66 - Your hostname, UDIC-GPU resolves to a loopback address: 127.0.0.1; using 134.193.130.129 instead (on interface enp4s0)
2018-07-10 22:45:24 WARN  Utils:66 - Set SPARK_LOCAL_IP if you need to bind to another address
2018-07-10 22:45:27 WARN  NativeCodeLoader:62 - Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
Welcome to

  ____
 /  __ \  _ __| | | |
/_ /  \ \| '_ \| | | |
 \___/\___/|_|_| |_|_|_|_| version 2.3.1
  /  \

Using Python version 2.7.12 (default, Dec  4 2017 14:50:18)
SparkSession available as 'spark'.
>>>
```

Create Dataframe in Pyspark

```
>>> e = sqlContext.createDataFrame([
...     ("a", "b", "friend"),
...     ("b", "c", "follow"),
...     ("c", "b", "follow"),
...     ("f", "c", "follow"),
...     ("e", "f", "follow"),
...     ("e", "d", "friend"),
...     ("d", "a", "friend"),
...     ("a", "e", "friend")
... ], ["src", "dst", "relationship"])
>>> g = GraphFrame(v, e)
>>> g.vertices.show()
+---+-----+---+
| id|  name|age|
+---+-----+---+
|  a|  Alice| 34|
|  b|   Bob| 36|
|  c|Charlie| 30|
|  d|  David| 29|
|  e| Esther| 32|
|  f|  Fanny| 36|
|  g|  Gabby| 60|
+---+-----+---+
```

Create Dataframe in Scala

```
// Vertex DataFrame
val v = sqlContext.createDataFrame(List(
  ("a", "Alice", 34),
  ("b", "Bob", 36),
  ("c", "Charlie", 30),
  ("d", "David", 29),
  ("e", "Esther", 32),
  ("f", "Fanny", 36),
  ("g", "Gabby", 60)
)).toDF("id", "name", "age")

// Edge DataFrame
val e = sqlContext.createDataFrame(List(
  ("a", "b", "friend"),
  ("b", "c", "follow"),
  ("c", "b", "follow"),
  ("f", "c", "follow"),
  ("e", "f", "follow"),
  ("e", "d", "friend"),
  ("d", "a", "friend"),
  ("a", "e", "friend")
)).toDF("src", "dst", "relationship")

// Create a GraphFrame
val g = GraphFrame(v, e)
```

GraphFrames Output


```
>>> g.vertices.show()
+----+-----+----+
| id|   name|age|
+----+-----+----+
|  a|  Alice| 34|
|  b|   Bob| 36|
|  c|Charlie| 30|
|  d| David| 29|
|  e| Esther| 32|
|  f|  Fanny| 36|
|  g|  Gabby| 60|
+----+-----+----+

>>> g.edges.show()
+----+-----+-----+
|src|dst|relationship|
+----+-----+-----+
|  a|  b|      friend|
|  b|  c|      follow|
|  c|  b|      follow|
|  f|  c|      follow|
|  e|  f|      follow|
|  e|  d|      friend|
|  d|  a|      friend|
|  a|  e|      friend|
+----+-----+-----+
```

Motif Finding


- Search for structural pattern in a graph

```
val paths: DataFrame =  
  g.find("(a)-[e1]->(b);  
        (b)-[e2]->(c);  
        !(c)-[]->(a)")
```




Continued ...

```
val paths: DataFrame =  
  g.find("(a) - [e1] -> (b);  
        (b) - [e2] -> (c);  
        !(c) - [] -> (a)")
```




Continued ...

```
val paths: DataFrame =
  g.find("(a) - [e1] -> (b);
        (b) - [e2] -> (c);
        ! (c) - [] -> (a)")
```



Continued ...

```
val paths: DataFrame =
  g.find("(a) - [e1] -> (b);
        (b) - [e2] -> (c);
        ! (c) - [] -> (a)")
```



Motifs filter by age (pyspark)

```
>>> numFollows = g.edges.filter("relationship = 'follow'").count()
>>> numFollows = g.edges.filter("relationship = 'follow'").count().show()
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
AttributeError: 'int' object has no attribute 'show'
>>> motifs = g.find("(a)-[e]->(b); (b)-[e2]->(a)")
>>> motifs.show()
+-----+-----+-----+-----+
|          a|          e|          b|          e2|
+-----+-----+-----+-----+
|[c, Charlie, 30]|[c, b, follow]|    [b, Bob, 36]|[b, c, follow]|
|    [b, Bob, 36]|[b, c, follow]|[c, Charlie, 30]|[c, b, follow]|
+-----+-----+-----+-----+

>>> motifs.filter("b.age > 30").show()
+-----+-----+-----+-----+
|          a|          e|          b|          e2|
+-----+-----+-----+-----+
|[c, Charlie, 30]|[c, b, follow]| [b, Bob, 36]|[b, c, follow]|
+-----+-----+-----+-----+
```

Scala

```
import org.graphframes.{examples, GraphFrame}
val g: GraphFrame = examples.Graphs.friends // get example graph

// Search for pairs of vertices with edges in both directions between them.
val motifs: GraphFrame = g.find("(a)-[e]->(b); (b)-[e2]->(a)")
motifs.show()

// More complex queries can be expressed by applying filters.
motifs.filter("b.age > 30").show()
```

Combination of Relationship (pyspark)

```
>>> from pyspark.sql.functions import col, lit, udf, when
>>> from pyspark.sql.types import IntegerType
>>> chain4 = g.find("(a)-[ab]->(b); (b)-[bc]->(c); (c)-[cd]->(d)")
>>>
>>> sumFriends = \
...     lambda cnt,relationship: when(relationship == "friend", cnt+1).otherwise(cnt)
>>> condition = \
...     reduce(lambda cnt,e: sumFriends(cnt, col(e).relationship), ["ab", "bc", "cd"], lit(0))
>>> chainWith2Friends2 = chain4.where(condition >= 2)
>>> chainWith2Friends2.show()
```

	a	ab	b	bc	c	cd	d
	[d, David, 29]	[d, a, friend]	[a, Alice, 34]	[a, e, friend]	[e, Esther, 32]	[e, f, follow]	[f, Fanny, 36]
	[e, Esther, 32]	[e, d, friend]	[d, David, 29]	[d, a, friend]	[a, Alice, 34]	[a, e, friend]	[e, Esther, 32]
	[d, David, 29]	[d, a, friend]	[a, Alice, 34]	[a, e, friend]	[e, Esther, 32]	[e, d, friend]	[d, David, 29]
	[d, David, 29]	[d, a, friend]	[a, Alice, 34]	[a, b, friend]	[b, Bob, 36]	[b, c, follow]	[c, Charlie, 30]
	[e, Esther, 32]	[e, d, friend]	[d, David, 29]	[d, a, friend]	[a, Alice, 34]	[a, b, friend]	[b, Bob, 36]
	[a, Alice, 34]	[a, e, friend]	[e, Esther, 32]	[e, d, friend]	[d, David, 29]	[d, a, friend]	[a, Alice, 34]

Scala

```
import org.apache.spark.sql.Column
import org.apache.spark.sql.functions.{col, when}
import org.graphframes.{examples, GraphFrame}

val g: GraphFrame = examples.Graphs.friends // get example graph

// Find chains of 4 vertices.
val chain4 = g.find("(a)-[ab]->(b); (b)-[bc]->(c); (c)-[cd]->(d)")

// Query on sequence, with state (cnt)
// (a) Define method for updating state given the next element of the motif.
def sumFriends(cnt: Column, relationship: Column): Column = {
  when(relationship === "friend", cnt + 1).otherwise(cnt)
}
// (b) Use sequence operation to apply method to sequence of elements in motif.
//     In this case, the elements are the 3 edges.
val condition = { Seq("ab", "bc", "cd")
  .foldLeft(lit(0))((cnt, e) => sumFriends(cnt, col(e)("relationship"))) }
// (c) Apply filter to DataFrame.
val chainWith2Friends2 = chain4.where(condition >= 2)
chainWith2Friends2.show()
```

Motif show

```
>>> motifs.show()
```

```
+-----+-----+-----+-----+
|           a|           e|           b|           e2|
+-----+-----+-----+-----+
|[c, Charlie, 30]|[c, b, follow]|    [b, Bob, 36]|[b, c, follow]|
|    [b, Bob, 36]|[b, c, follow]|[c, Charlie, 30]|[c, b, follow]|
+-----+-----+-----+-----+
```


Create Subgraphs (Pyspark)

```
from graphframes.examples import Graphs
g = Graphs(sqlContext).friends() # Get example graph

# Select subgraph of users older than 30, and edges of type "friend"
v2 = g.vertices.filter("age > 30")
e2 = g.edges.filter("relationship = 'friend'")
g2 = GraphFrame(v2, e2)
```

Create Subgraphs (Scala)


```
import org.graphframes.examples
val g: GraphFrame = examples.Graphs.friends

// Select subgraph of users older than 30, and edges of type "friend"
val v2 = g.vertices.filter("age > 30")
val e2 = g.edges.filter("relationship = 'friend'")
val g2 = GraphFrame(v2, e2)
```

Graph Algorithms

Breadth-first-search (BFS) algorithm

- Breadth-first search (BFS) finds the shortest path(s) from one vertex (or a set of vertices) to another vertex (or a set of vertices)
- The beginning and end vertices are specified as Spark DataFrame expressions



```

val g: GraphFrame = examples.Graphs.friends // get example graph

// Search from "Esther" for users of age <= 32.
val paths: DataFrame = g.bfs.fromExpr("name = 'Esther']").toExpr("age < 32").run()
paths.show()

// Specify edge filters or max path lengths.
g.bfs.fromExpr("name = 'Esther']").toExpr("age < 32")
  .edgeFilter("relationship != 'friend'")
  .maxPathLength(3)
  .run()

```



```

g = Graphs(sqlContext).friends() # Get example graph

# Search from "Esther" for users of age < 32.
paths = g.bfs("name = 'Esther'", "age < 32")
paths.show()


# Specify edge filters or max path lengths.
g.bfs("name = 'Esther'", "age < 32", \
      edgeFilter="relationship != 'friend'", maxPathLength=3)

```



Connected Components

- Computes the connected component membership of each vertex and returns a graph with each vertex assigned a component ID.



```
val g: GraphFrame = examples.Graphs.friends // get example graph
```

```
val result = g.connectedComponents.run()
```

```
result.select("id", "component").orderBy("component").show()
```



```
g = Graphs(sqlContext).friends() # Get example graph
```


```
result = g.connectedComponents()
```

```
result.select("id", "component").orderBy("component").show()
```



Strongly Connected Components

- Compute the strongly connected component (SCC) of each vertex and return a graph with each vertex assigned to the SCC containing that vertex.




```
val g: GraphFrame = examples.Graphs.friends // get example graph

val result = g.stronglyConnectedComponents.maxIter(10).run()
result.select("id", "component").orderBy("component").show()
```




```
g = Graphs(sqlContext).friends() # Get example graph

result = g.stronglyConnectedComponents(maxIter=10)
result.select("id", "component").orderBy("component").show()
```



PageRank



- PageRank (PR) is an algorithm used by Google Search to rank websites in their search engine results. PageRank was named after Larry Page
- PageRank works by counting the number and quality of links to a page to determine a rough estimate of how important is the website

```

val g: GraphFrame = examples.Graphs.friends // get example graph

// Run PageRank until convergence to tolerance "tol".
val results = g.pageRank.resetProbability(0.15).tol(0.01).run()
// Display resulting pageranks and final edge weights
// Note that the displayed pagerank may be truncated, e.g., missing the E notation.
// In Spark 1.5+, you can use show(truncate=false) to avoid truncation.
results.vertices.select("id", "pagerank").show()
results.edges.select("src", "dst", "weight").show()

// Run PageRank for a fixed number of iterations.
val results2 = g.pageRank.resetProbability(0.15).maxIter(10).run()

// Run PageRank personalized for vertex "a"
val results3 = g.pageRank.resetProbability(0.15).maxIter(10).sourceId("a").run()

```



```

g = Graphs(sqlContext).friends() # Get example graph

# Run PageRank until convergence to tolerance "tol".
results = g.pageRank(resetProbability=0.15, tol=0.01)
# Display resulting pageranks and final edge weights
# Note that the displayed pagerank may be truncated, e.g., missing the E notation.
# In Spark 1.5+, you can use show(truncate=False) to avoid truncation.
results.vertices.select("id", "pagerank").show()
results.edges.select("src", "dst", "weight").show()


# Run PageRank for a fixed number of iterations.
results2 = g.pageRank(resetProbability=0.15, maxIter=10)

# Run PageRank personalized for vertex "a"
results3 = g.pageRank(resetProbability=0.15, maxIter=10, sourceId="a")

```



Triangle Counting



Computes the number of triangles passing through each vertex.

```
val g: GraphFrame = examples.Graphs.friends // get example graph

val results = g.triangleCount.run()
results.select("id", "count").show()
```



```
g = Graphs(sqlContext).friends() # Get example graph

results = g.triangleCount()
results.select("id", "count").show()
```



Save and Loading Graphframe

```
val g: GraphFrame = examples.Graphs.friends // get example graph

// Save vertices and edges as Parquet to some location.
g.vertices.write.parquet("hdfs://myLocation/vertices")
g.edges.write.parquet("hdfs://myLocation/edges")

// Load the vertices and edges back.
val sameV = sqlContext.read.parquet("hdfs://myLocation/vertices")
val sameE = sqlContext.read.parquet("hdfs://myLocation/edges")

// Create an identical GraphFrame.
val sameG = GraphFrame(sameV, sameE)
```



```
g = Graphs(sqlContext).friends() # Get example graph

# Save vertices and edges as Parquet to some location.
g.vertices.write.parquet("hdfs://myLocation/vertices")
g.edges.write.parquet("hdfs://myLocation/edges")

# Load the vertices and edges back.
sameV = sqlContext.read.parquet("hdfs://myLocation/vertices")
sameE = sqlContext.read.parquet("hdfs://myLocation/edges")

# Create an identical GraphFrame.
sameG = GraphFrame(sameV, sameE)
```



References

- <https://spark.apache.org/docs/latest/graphx-programming-guide.html>
- <http://www.sparktutorials.net/Analyzing+Flight+Data%3A+A+Gentle+Introduction+to+GraphX+in+Spark>
<http://www.sparktutorials.net/Analyzing+Flight+Data%3A+A+Gentle+Introduction+to+GraphX+in+Spark>
- <https://mapr.com/blog/how-get-started-using-apache-spark-graphx-scala/>
- <https://www.edureka.co/blog/spark-graphx/>
- <https://graphframes.github.io/user-guide.html>
- <https://databricks.com/blog/2016/03/03/introducing-graphframes.html>
- <https://docs.databricks.com/spark/latest/graph-analysis/graphframes/user-guide-scala.html>